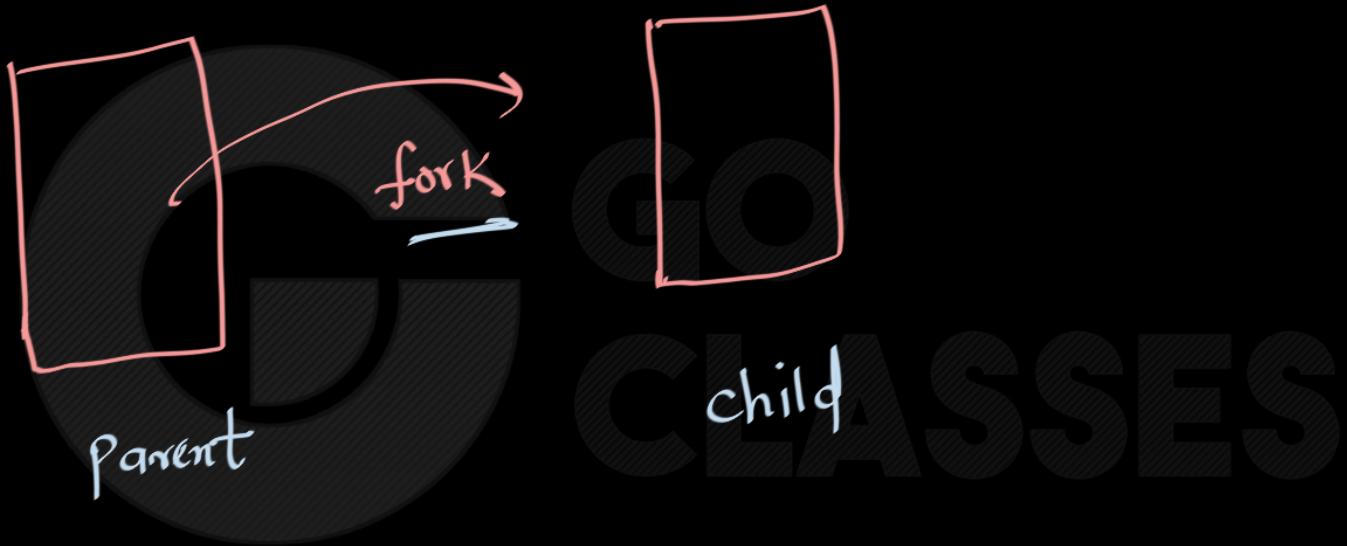




G GO Lecture 4 CLASSES



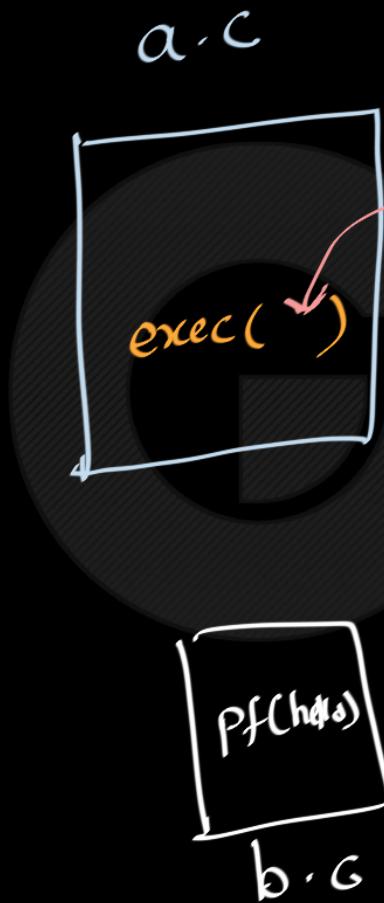


Operating Systems



exec System call





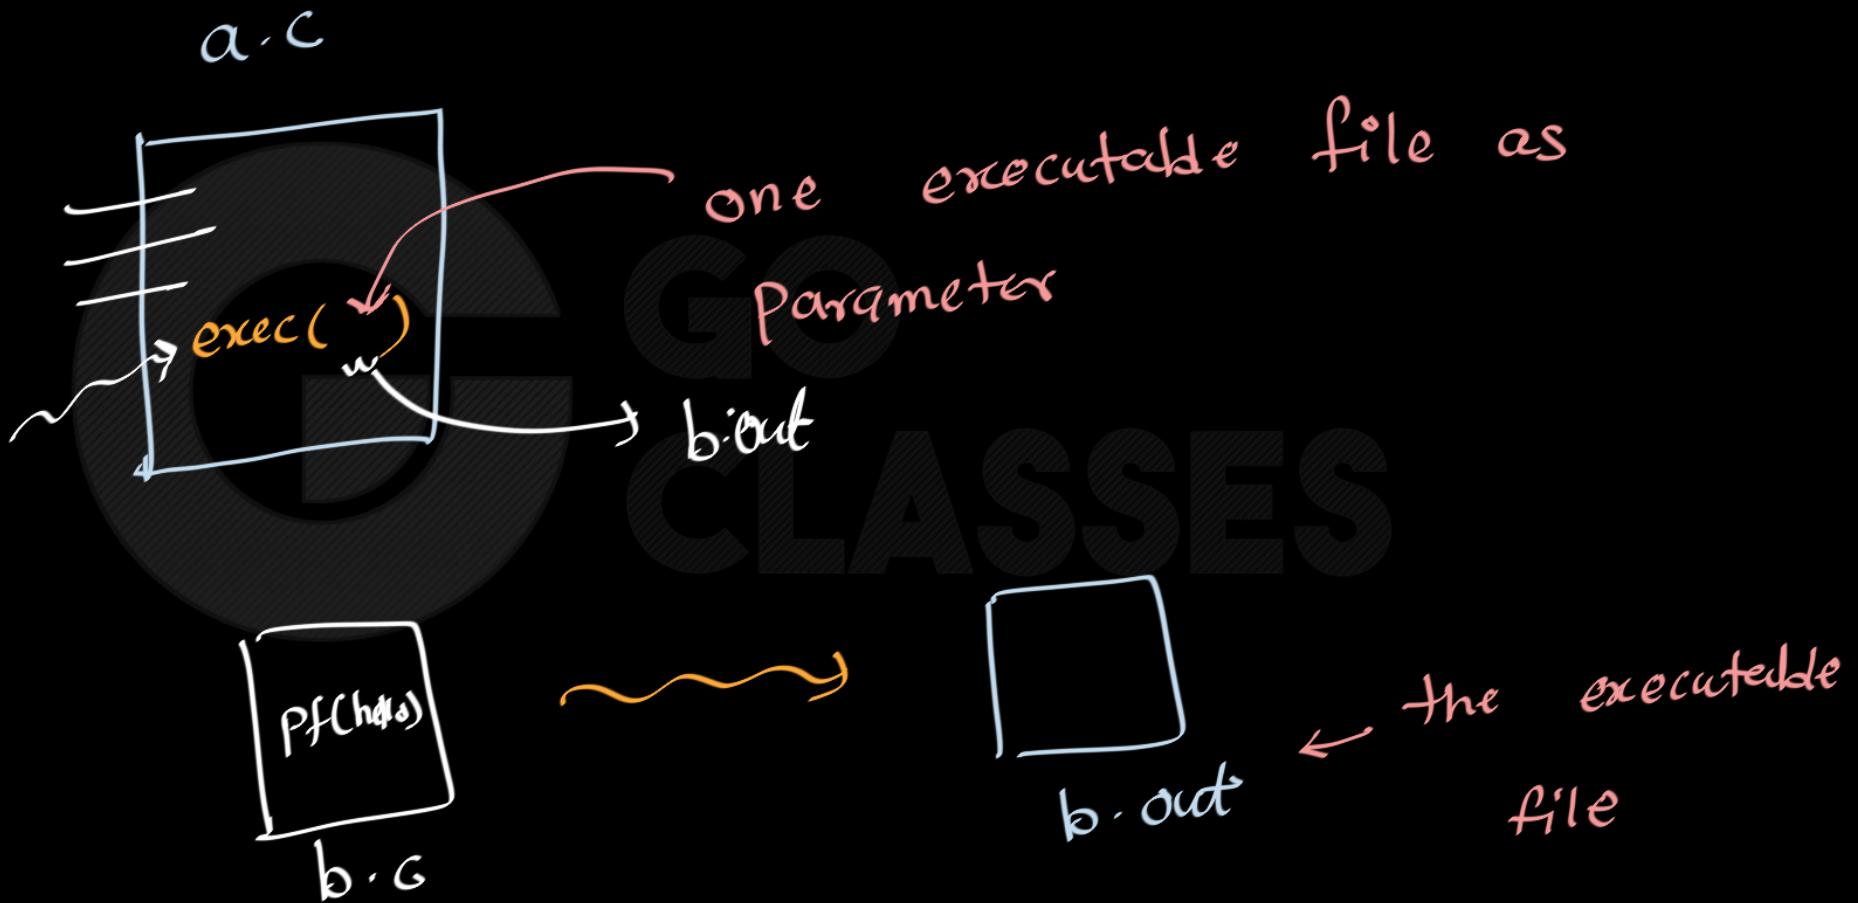
one executable file as
Parameter

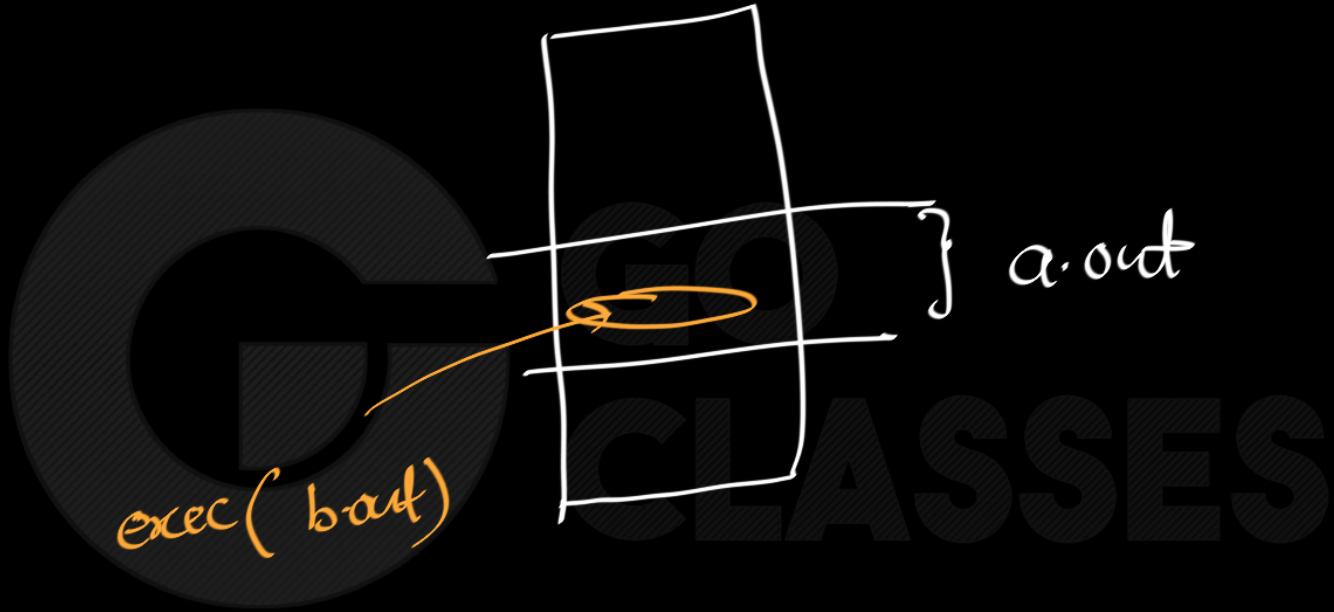
~~~~~

b.out

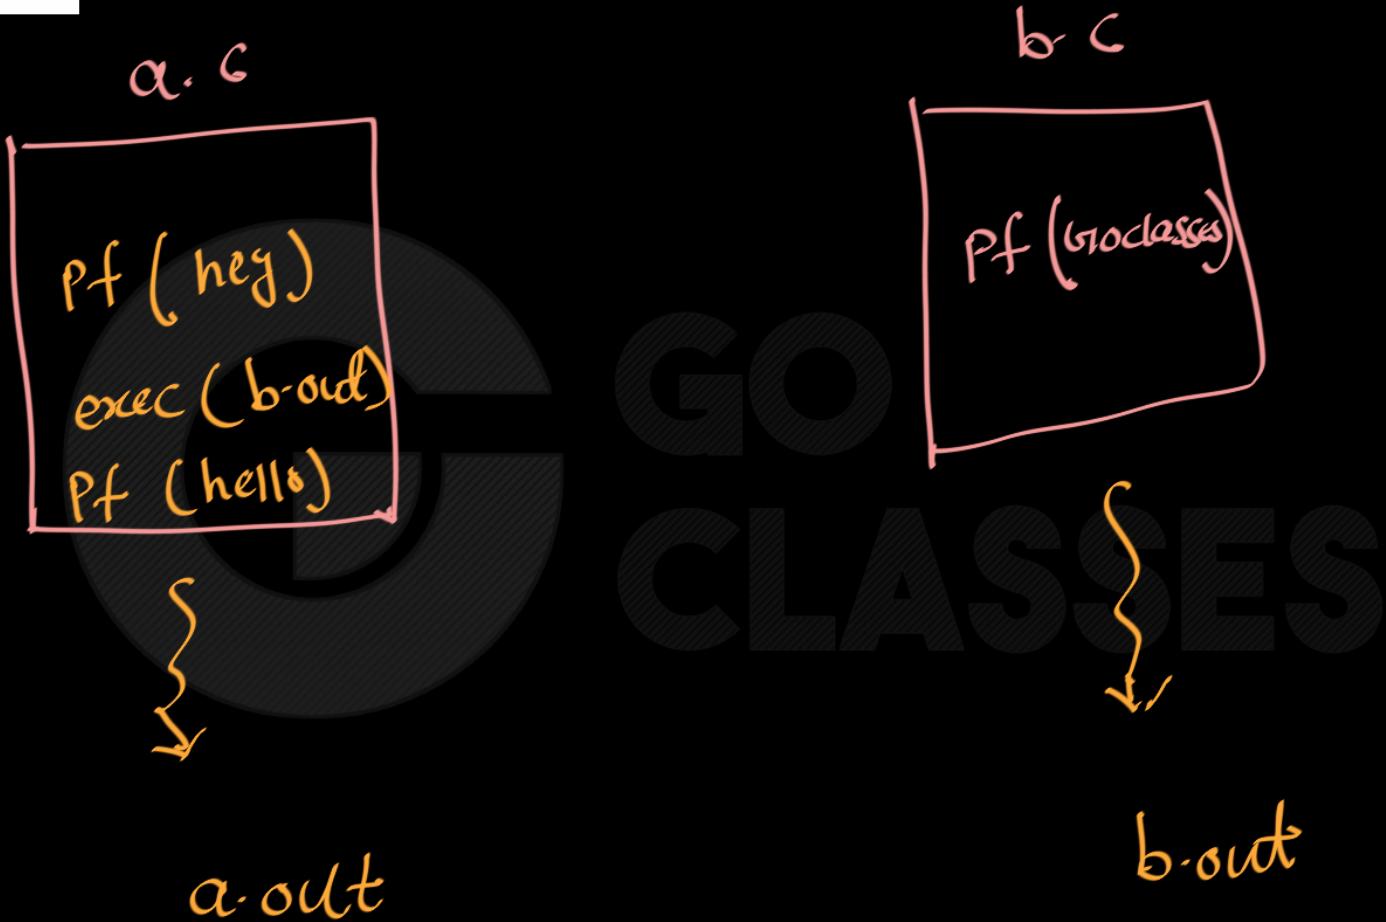
the executed file

The text "one executable file as Parameter" is written above the "p.f(hello)" box. Below it, a wavy arrow points to a third box labeled "b.out". To the right of "b.out", the text "the executed file" is written with an arrow pointing towards it.





a.out will get replaced with b.out



question :

what is the output  
we get if we run  
a.out file?

a.out

pf (hey)  
exec (b.out)  
pf (hello)



a.out

b.out

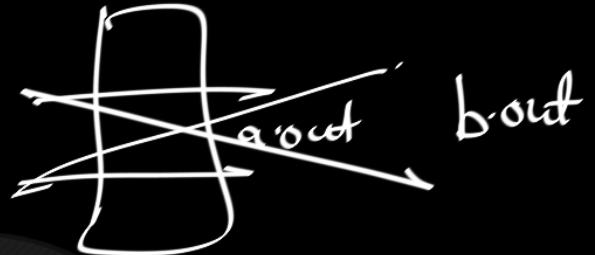
pf (b.out)



b.out

question :

what is the output  
we get if we run  
a.out file?



a. c

pf (neg)  
exec (b.out)  
pf (hello)



a.out

b.c

pf (bioclasses)



b.out

Answer

hey bioclasses

The screenshot shows a terminal window with two tabs: 'a' and 'b'. The 'a' tab contains the following C code:

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main() {
5     printf("hey its a.c\n");
6     fflush(stdout); ← ignore
7
8     char *char_array[] = { "Welcome", "To", "GOClasses", NULL };
9     execv("b.out", char_array);
10
11    printf("Welcome back to a.c");
12
13    return 0;
14 }
15
```

The 'b' tab contains the following C code:

```
b.c } { int main() {
printf("hey its b.c");
return 0;
}
```



The image shows a terminal window with two code editors. The left editor is titled 'c a' and contains the following C code:

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main() {
5     printf("hey its a.c\n");
6     fflush(stdout);
7
8     char *char_array[] = { "Welcome", "To",
9         "GOClasses", NULL };
10    execv("b.out", char_array);
11
12    printf("Welcome back to a.c");
13
14    return 0;
15 }
```

The right editor is titled 'c b' and contains the following C code:

```
1 #include <stdio.h>
2
3
4 int main() {
5     printf("hey its b.c");
6     return 0;
7
8
9 }
```

Below the editors is a terminal window titled 'OSCodes -- zsh -- 80x24'. It shows the output of running the program:

```
(base) sachinmittal@Sachins-MacBook-Pro OSCodes % ./a.out
hey its a.c
hey its b.c
(base) sachinmittal@Sachins-MacBook-Pro OSCodes %
```



# Operating Systems

The screenshot shows a terminal window with a dark background. At the top, there are three colored dots (yellow, green, blue) and a file icon labeled 'b.c'. Below the title bar, there are navigation icons: back, forward, and a search bar containing 'c b'. A message 'No Selection' is displayed. The main area contains the following C code:

```
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <stdlib.h>
4
5 int main() {
6     printf("hey its b.c");
7     return 0;
8 }
9
```



After Running a.c we get ↪

hey its a.c

hey its b.c

Note that we will not get Welcome back to a.c



Next Topic:  
Process Creation



Objective :

understand how does the 2 steps get

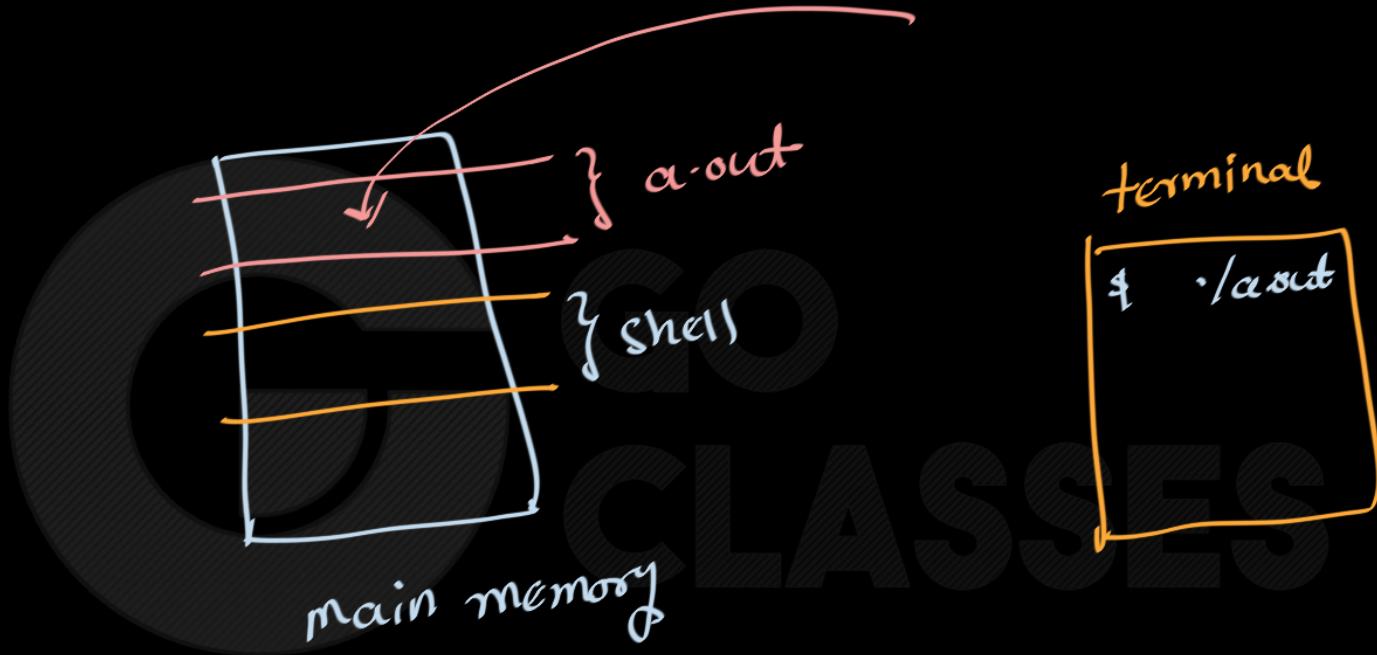
perform.

fork exec {

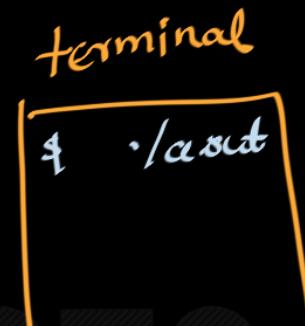
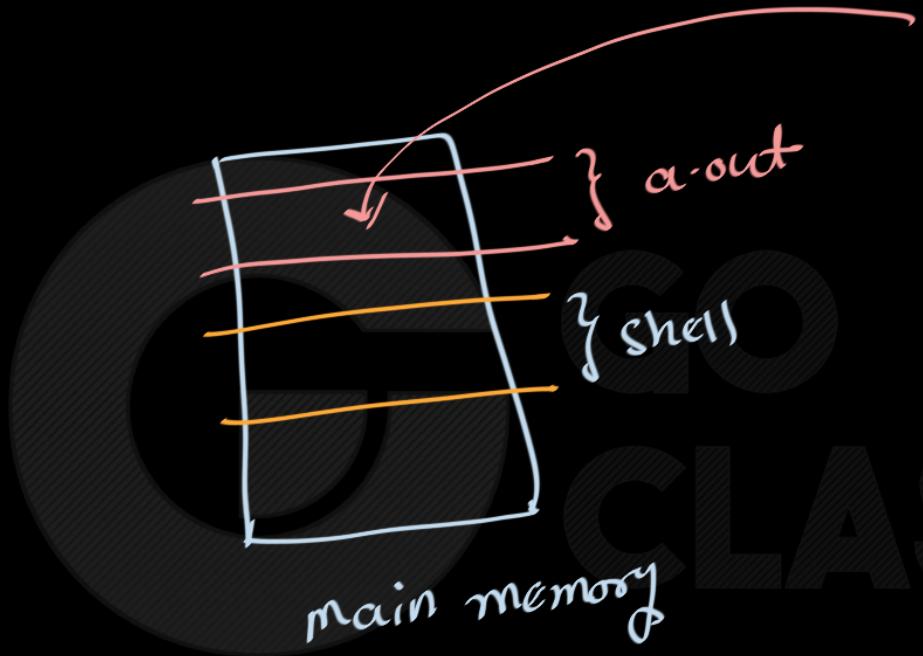
1. create space for hello.out
2. load hello.out in that space



- understand fork
- understand exec
- we will come back to process creation.



- ⇒ create Space for a.out
- ⇒ load the a.out in that space.



- ⇒ create Space for a.out
- ⇒ load the a.out in that space.

Sangeet Bhowmick to Everyone 8:13 PM



first fork will create a replica of parent process and exec will replace that process

Satyam Naik to Everyone 8:14 PM



fork exec



Running the command `ls` in a shell:

## Executing a new program

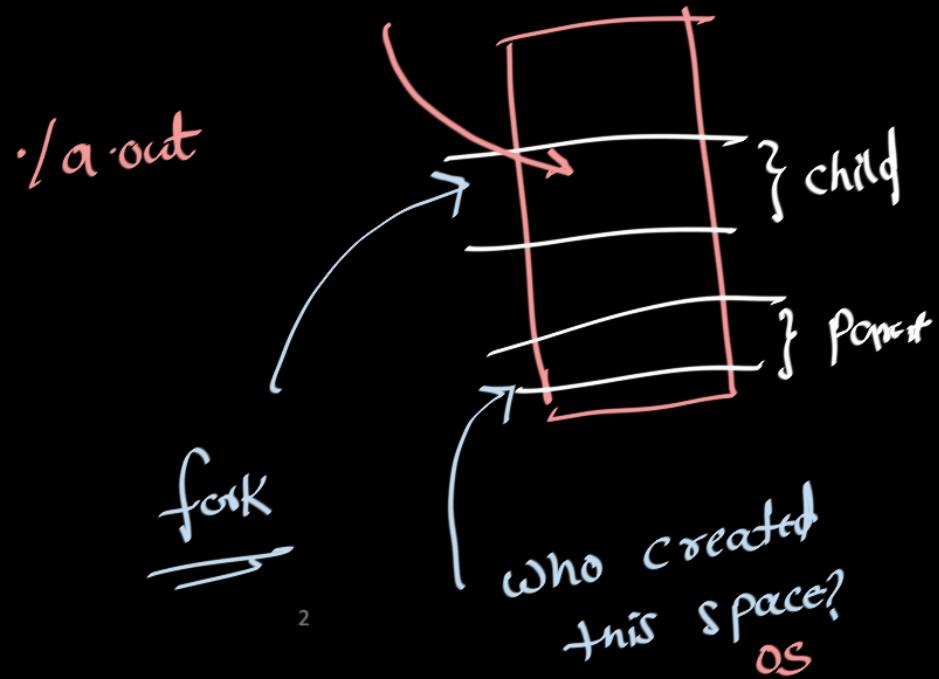
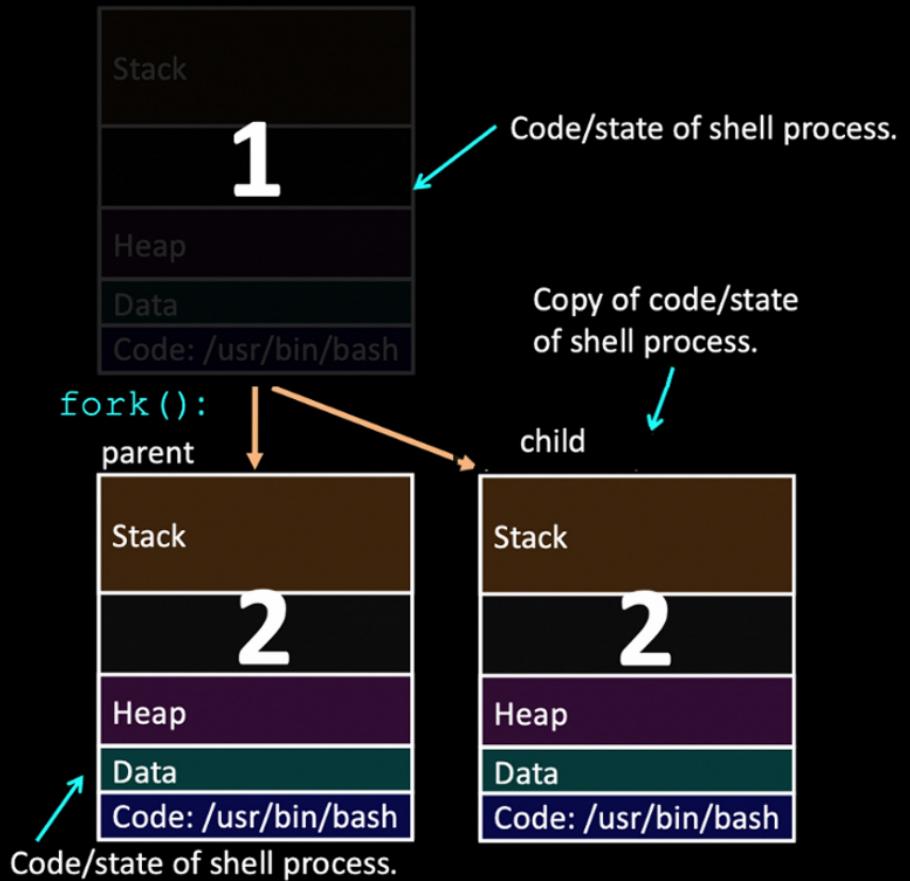


Code/state of shell process.



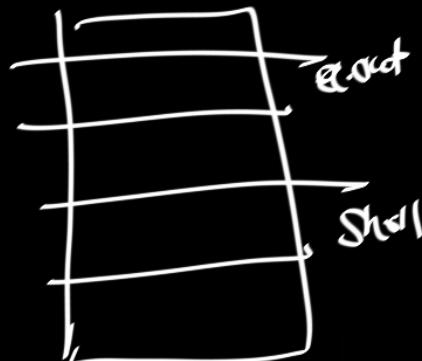
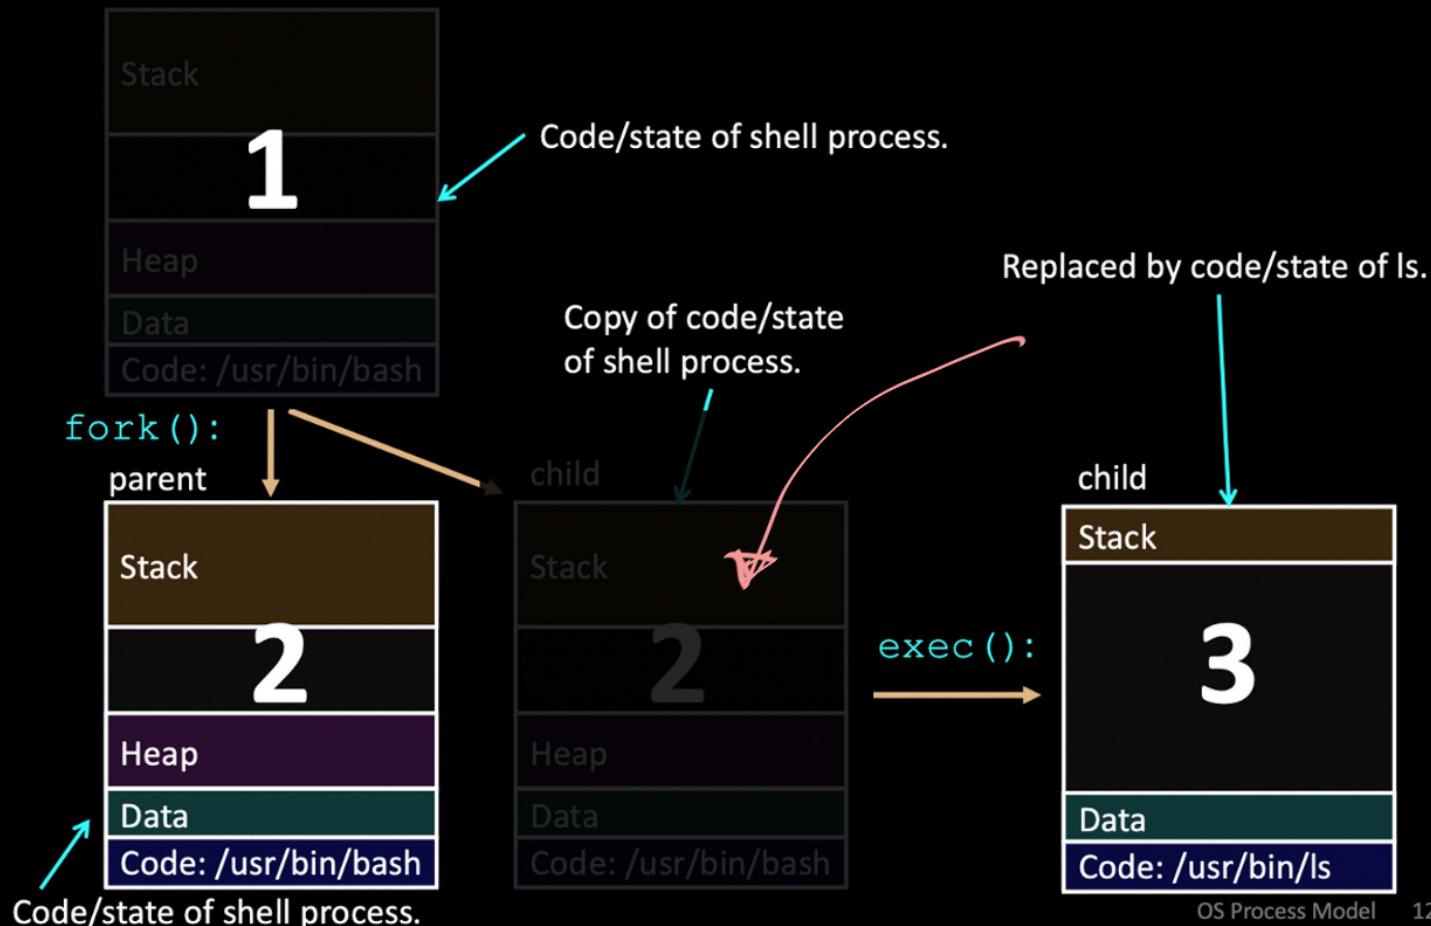
Running the command `ls` in a shell:

## Executing a new program



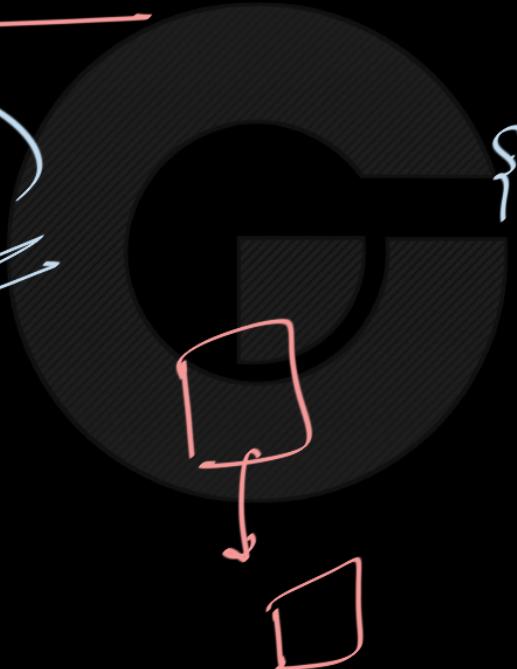
Running the command `ls` in a shell:

## Executing a new program



Shell code -

(optional)



```
while (true)
{
    scanf ("%s", &t);
    pid = fork();
    if (pid == 0)
        exec (t);
    else
        wait ( );
}
```

Some command



# Question

You write a UNIX shell, but instead of calling fork() then exec() to launch a new job, you instead insert a subtle difference: the code first calls exec() and then calls fork() like the following:

```
shell (...) {  
    ...  
    exec (cmd, args);  
    fork();  
    ...  
}
```

b-out

↙

ES

Does it work? What is the impact of this change to the shell, if any? (Explain)

means just like earlier

[https://www.eecg.toronto.edu/~yuan/teaching/ece344/pastExams/2013midterm\\_solution.pdf](https://www.eecg.toronto.edu/~yuan/teaching/ece344/pastExams/2013midterm_solution.pdf)

<https://www.cs.princeton.edu/courses/archive/fall13/cos318/exams/midterm-solutions-2013.pdf>



# Operating Systems

## Question

You write a UNIX shell, but instead of calling fork() then exec() to launch a new job, you instead insert a subtle difference: the code first calls exec() and then calls fork() like the following:

```
shell (...) {  
    ...  
    exec (cmd, args);  
    fork();  
    ...  
}
```

Answer: No

Does it work? What is the impact of this change to the shell, if any? (Explain)

→ we are destroying the shell

→ we are just replacing shell with bout

→ fork() will never get executed

[https://www.eecg.toronto.edu/~yuan/teaching/ece344/pastExams/2013midterm\\_solution.pdf](https://www.eecg.toronto.edu/~yuan/teaching/ece344/pastExams/2013midterm_solution.pdf)

<https://www.cs.princeton.edu/courses/archive/fall13/cos318/exams/midterm-solutions-2013.pdf>



## Question 11

**Part (a) [3 MARKS]** whatA

```
int main() {  
    int i = 0;  
    printf("%d ", i);  
    i = i + 1;  
    if (i >= 2)  
        exit(0);  
    execv("./whatA");  
}
```

D

- 
- A: 0 1 2
  - B: 0 0 0
  - C: 0 0 0 0 0 ... (forever)
  - D: 0 1 2 0 1 2 ... (forever)
  - E: none of the above
-

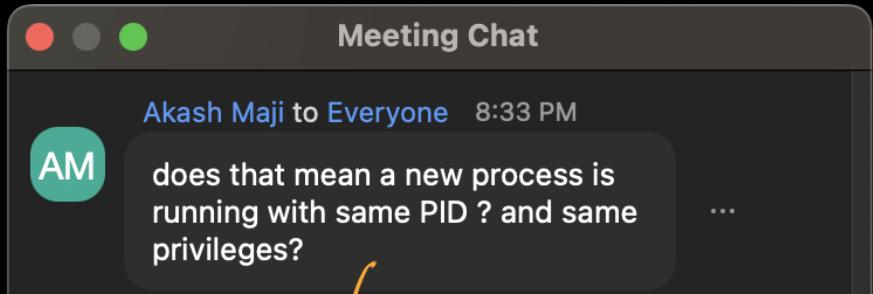


## Part (a) [3 MARKS] whatA

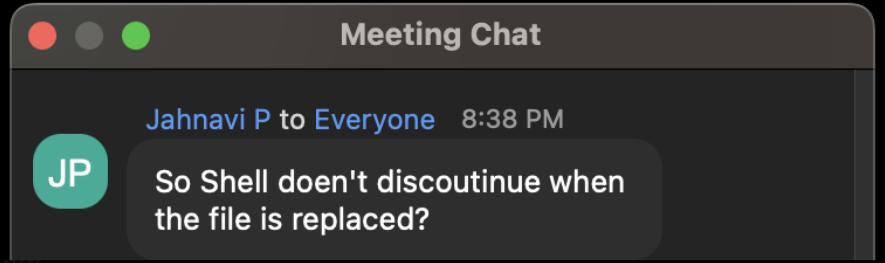
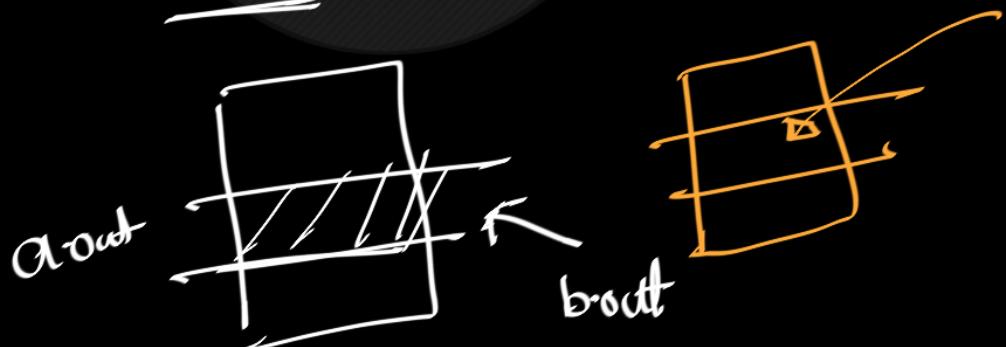
```
int main() {  
    int i = 0;  
    char *argv[2] = { "./whatA", NULL };  
  
    printf("%d ", i);  
    i = i + 1;  
    if (i >= 2)  
        exit(0);  
    execv(argv[0], argv);  
}
```

SES

- 
- A: 0 1 2
  - B: 0 0 0
  - C: 0 0 0 0 0 ... (forever)
  - D: 0 1 2 0 1 2 ... (forever)
  - E: none of the above



Yes PID is same



block state → } shell  
then it will  
not even get  
scheduled  
unless b.out  
finish b.out  
Main memory



# Operating Systems

## System Calls

- > what is system call ?
- > why they are always in lime light ?



## User Trust Issue

- Most of the users are not experts in computers and do not understand how hardware works. Thus they may not know how to properly manage the resources.
- An user (expert or not) can make incorrect operations while managing the resources (like handling h/w)
- An user can be malicious i.e. he/she can deliberately make incorrect operations. This can lead to system failure.

Solution from OS side: Never trust a user while dealing with critical tasks

$$\left[ \begin{array}{l} \text{int } a, b, c = 0 \\ a = b + c \end{array} \right]$$

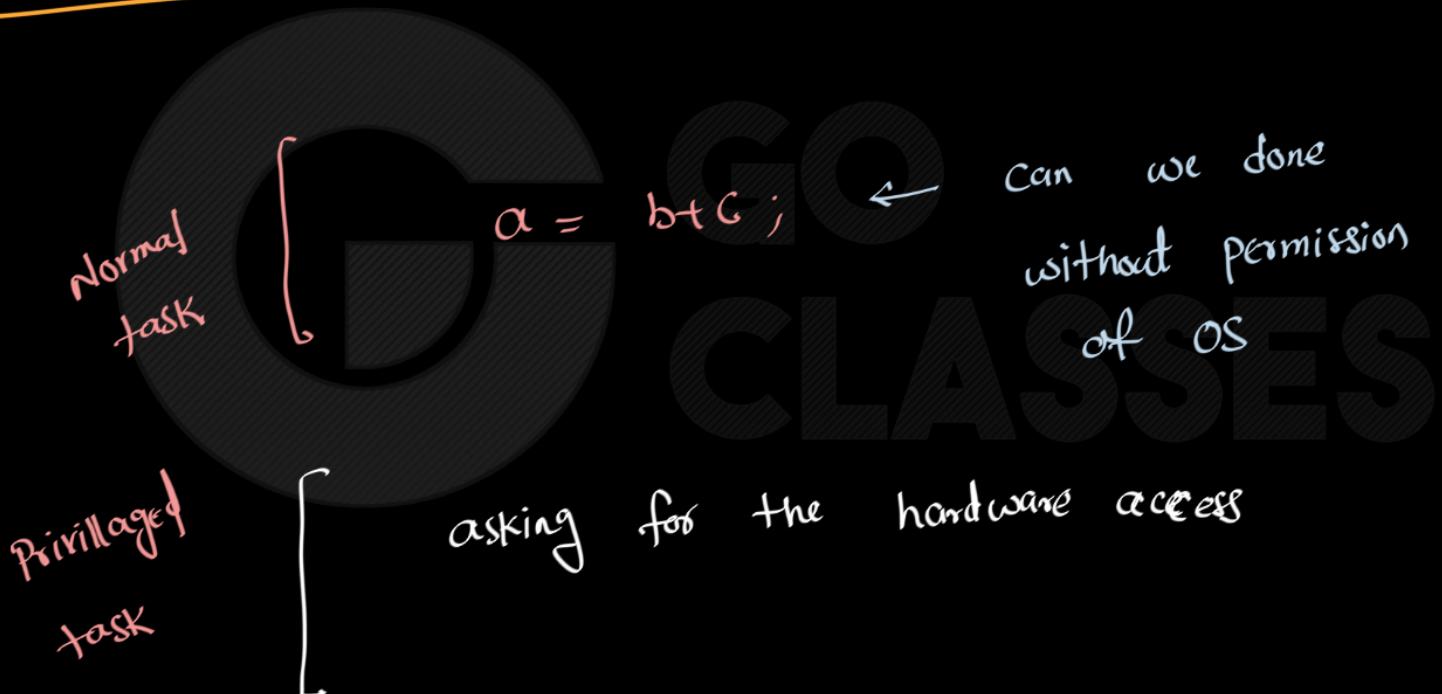
Normal  
task

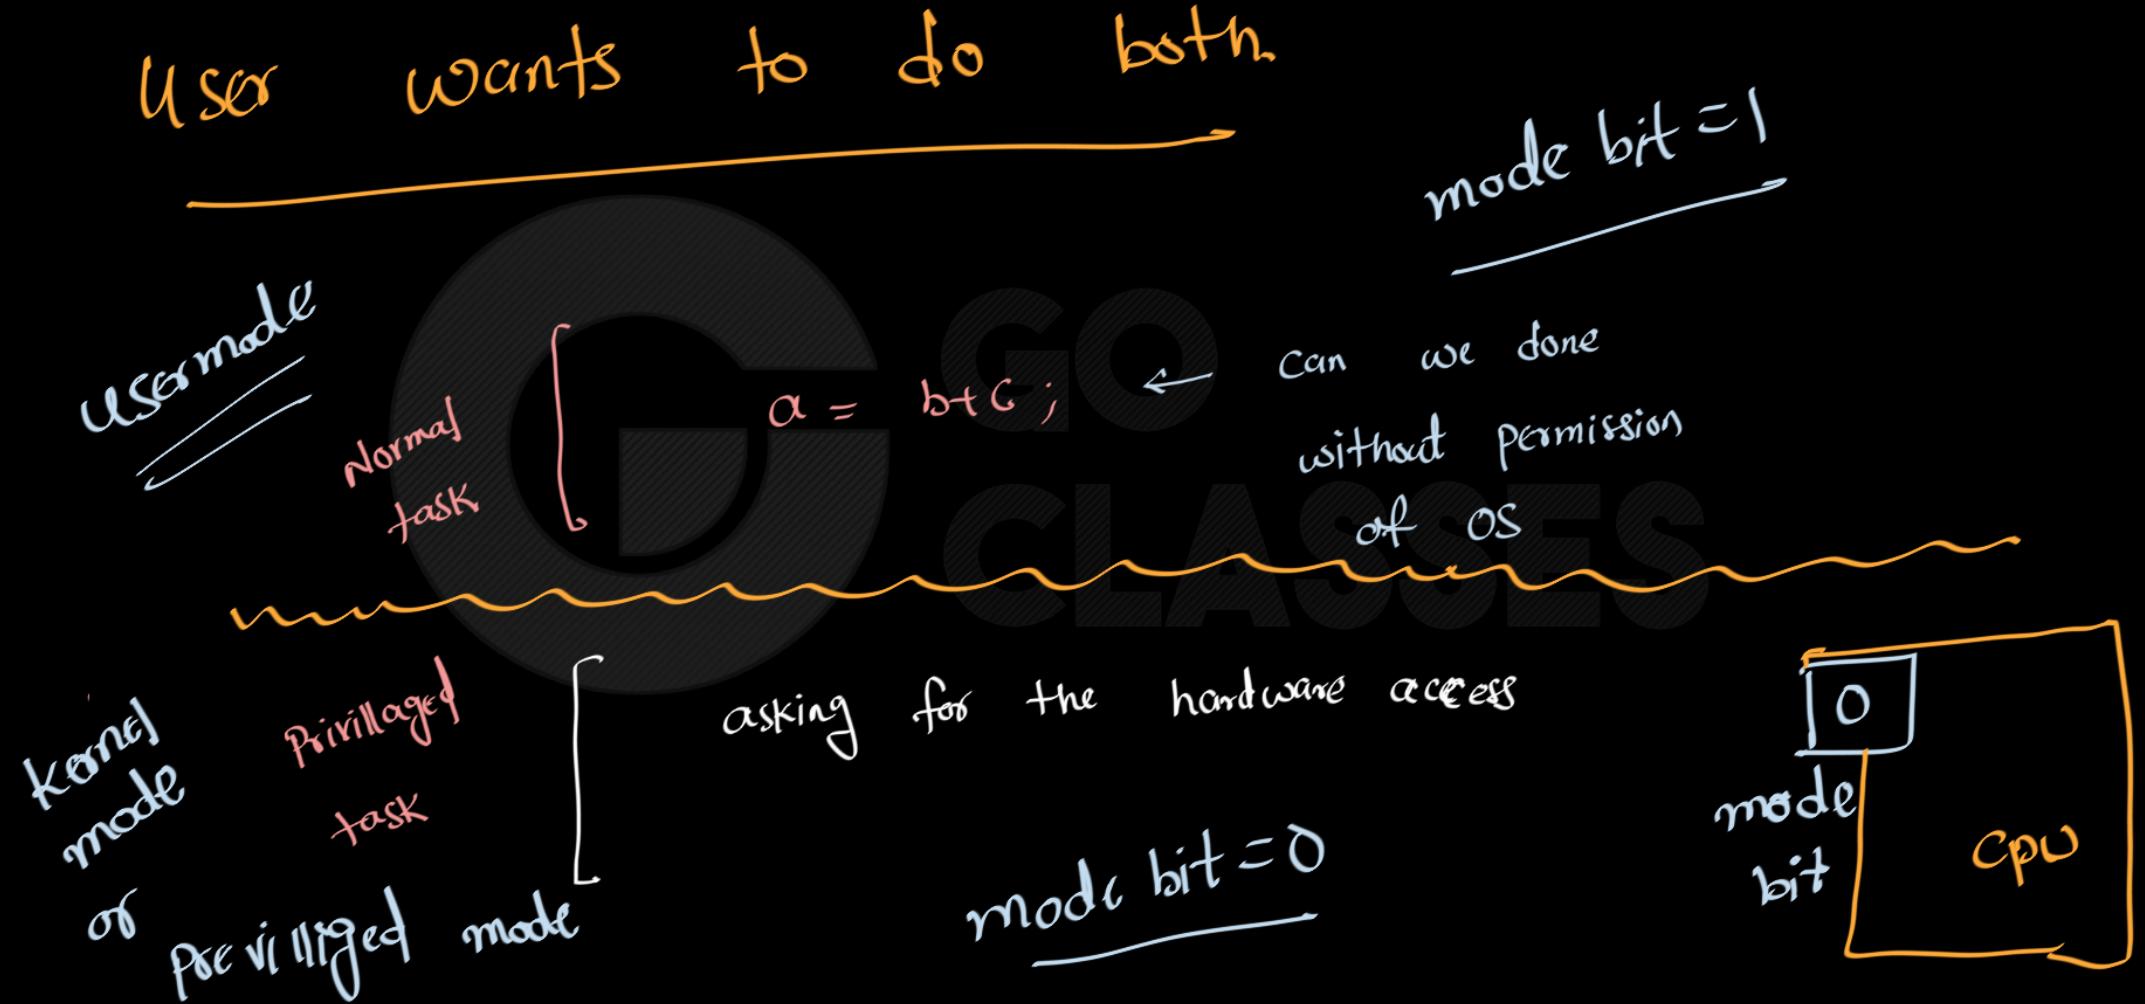
Privileged  
task

$a = b + c;$  ← Can we do  
without permission  
of OS

[ asking for the hardware access ]

User wants to do both.







## How user ask for OS Services ?

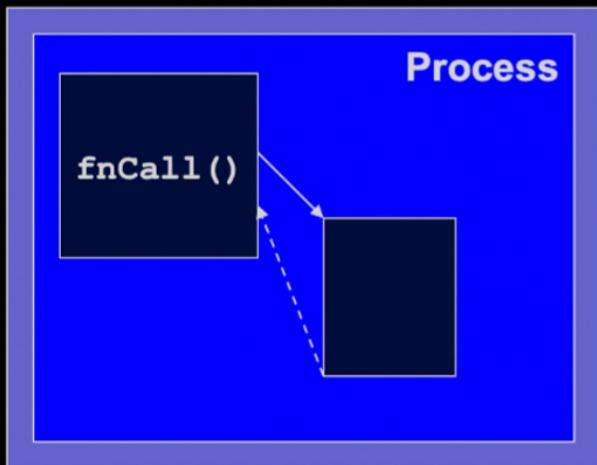
- Using System call

entry points to the kernel



# System Calls versus Function Calls

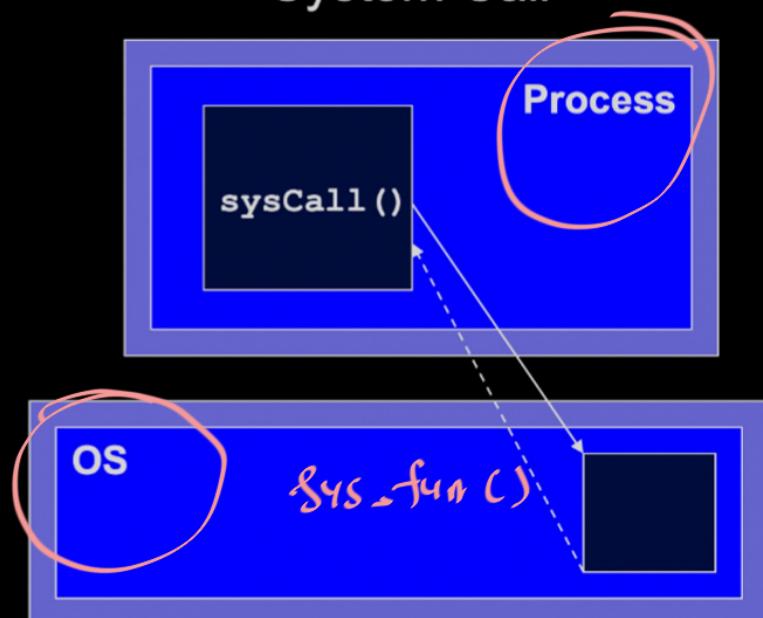
Function Call



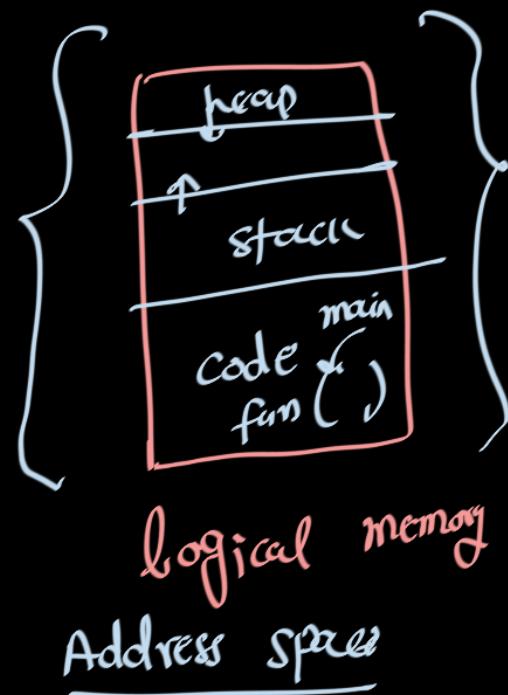
Caller and callee are in the same Process

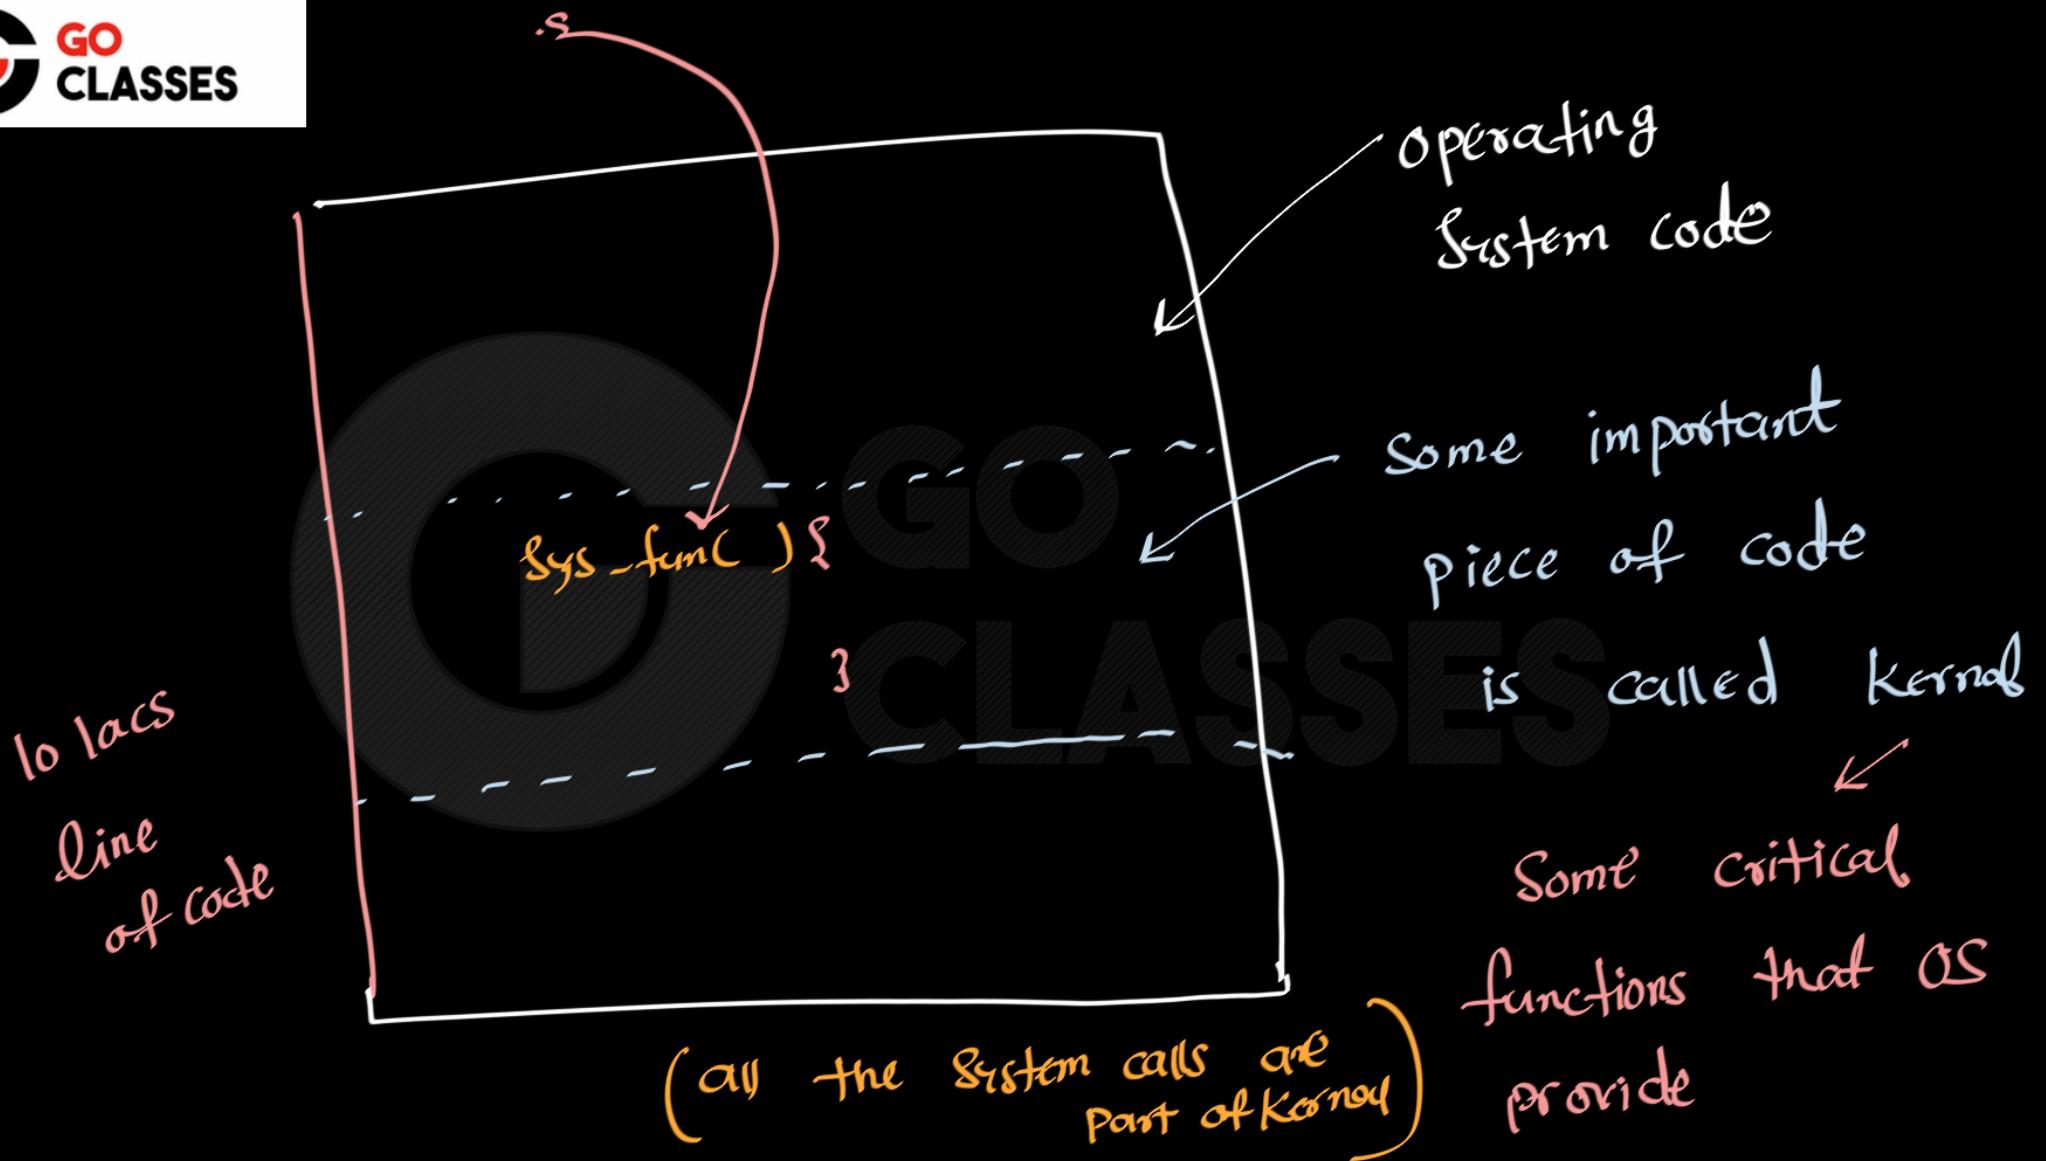
- Same user
- Same “domain of trust”

System Call

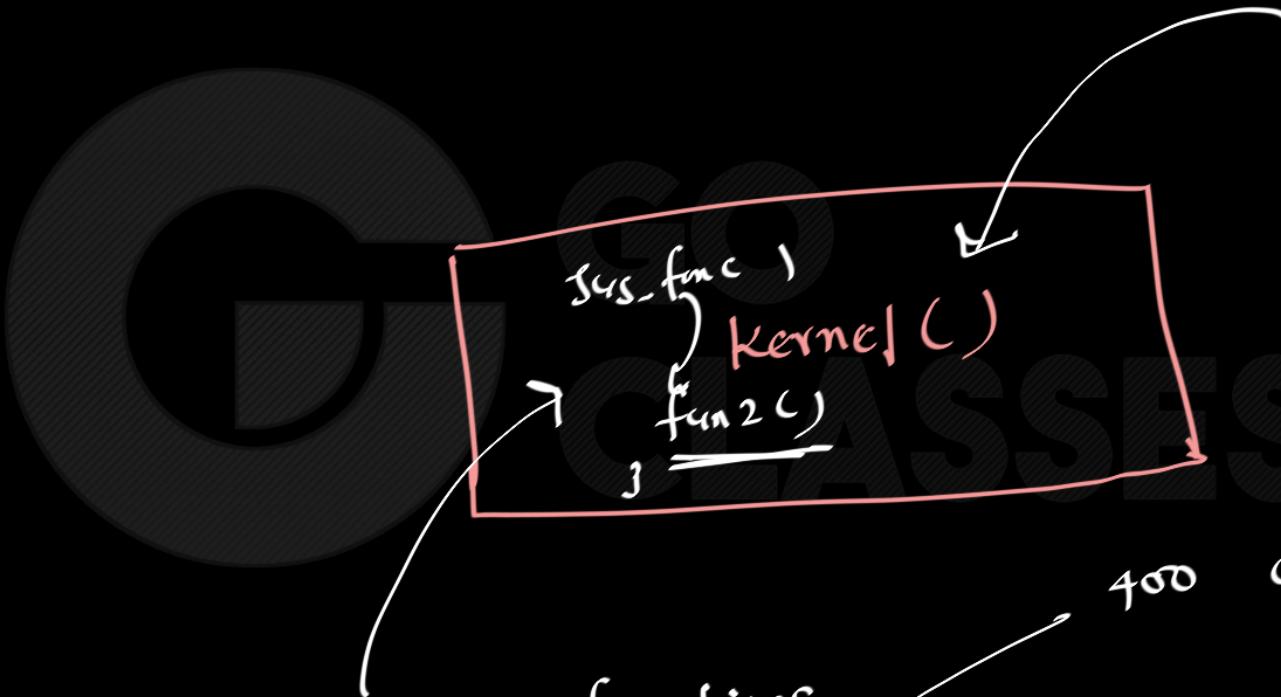


- OS is trusted; user is not.
- OS has super-privileges; user does not
- Must take measures to prevent abuse





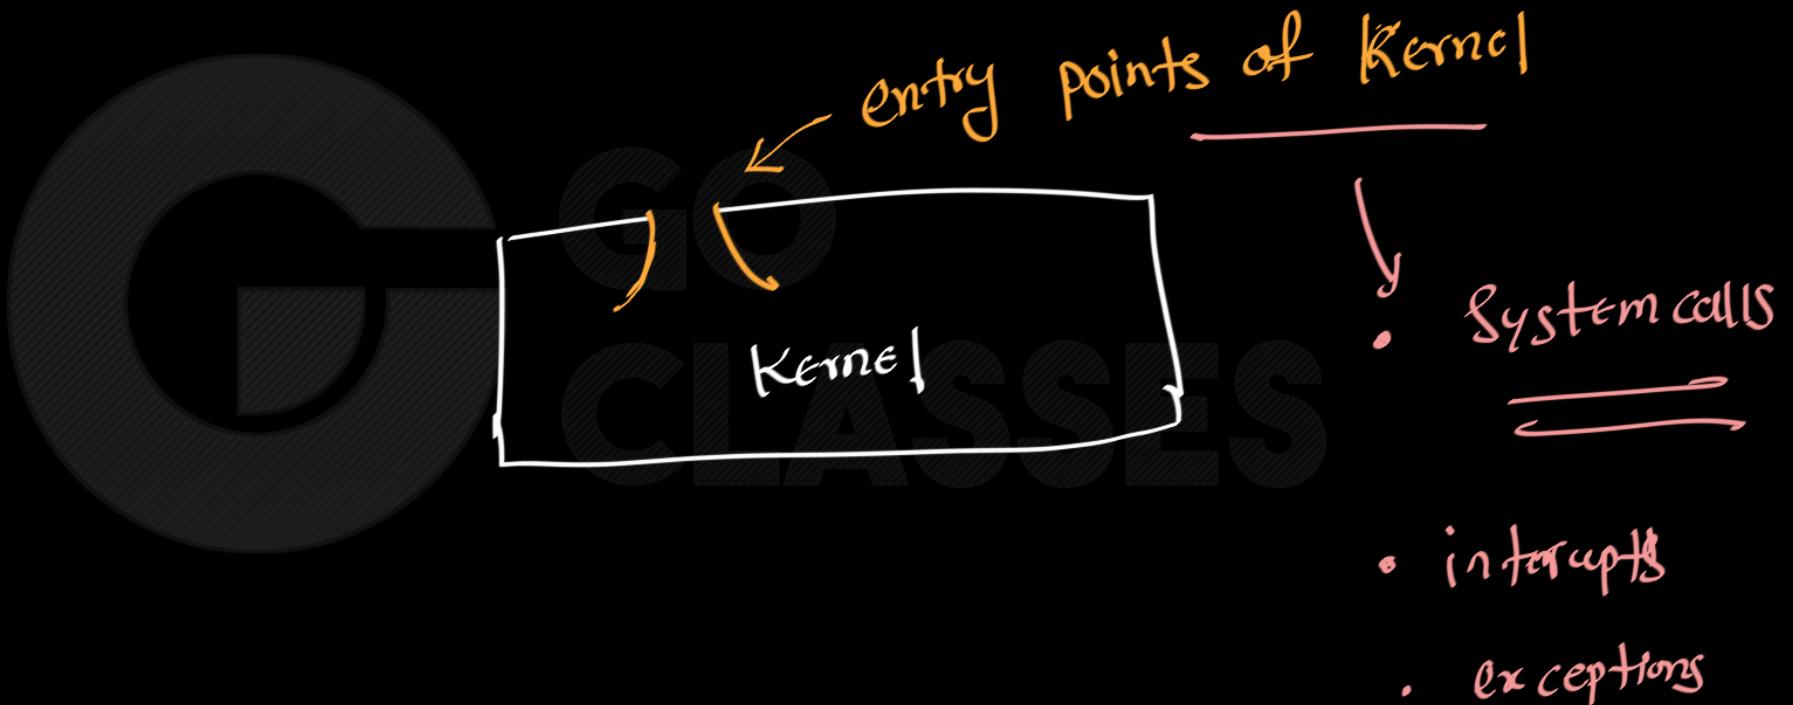
user code



1000 functions

400 are system calls

600 are not system calls





# Dual Mode of Operations

- User Mode

- Compute operations like addition, subtraction, etc.
- Reading and writing into program's memory

- Kernel Mode

- CPU and memory management
- I/O handling

How to ensure that user does not perform some thing illegal? Before running each instruction check whether user is allowed to run it or not. This check can be done efficiently in hardware.



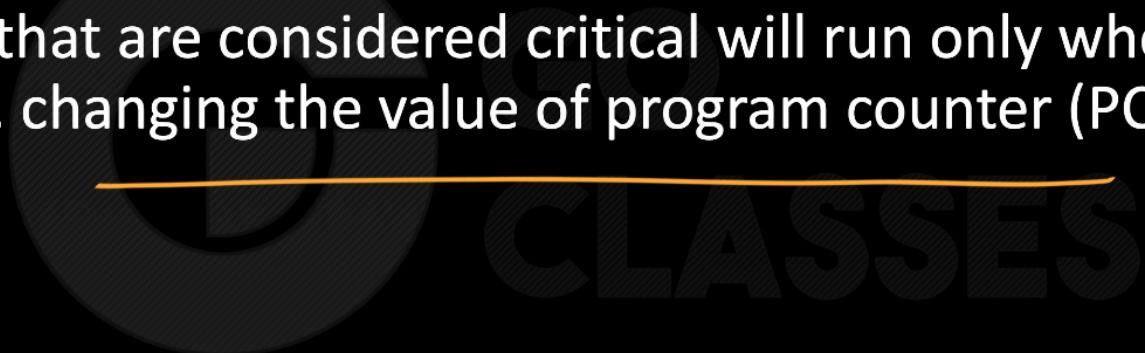
## Processor Modes

- The OS must restrict what a user process can do
  - What instructions can execute
  - What portions of the address space are accessible
- Supervisor mode (or kernel mode)
  - Can execute any instructions in the instruction set
    - Including halting the processor, changing mode bit, initiating I/O
  - Can access any memory location in the system
    - Including code and data in the OS address space
- User mode
  - Restricted capabilities
    - Cannot execute privileged instructions
    - Cannot directly reference code or data in OS address space
  - Any such attempt results in a fatal “protection fault”
    - Instead, access OS code and data indirectly via system calls



# Mode Bit and Privileged Instructions

- CPU maintains a special mode bit. If mode = 0 then system is in kernel/privileged mode; otherwise in user mode
- Instructions that are considered critical will run only when mode = 0. For example, changing the value of program counter (PC). 



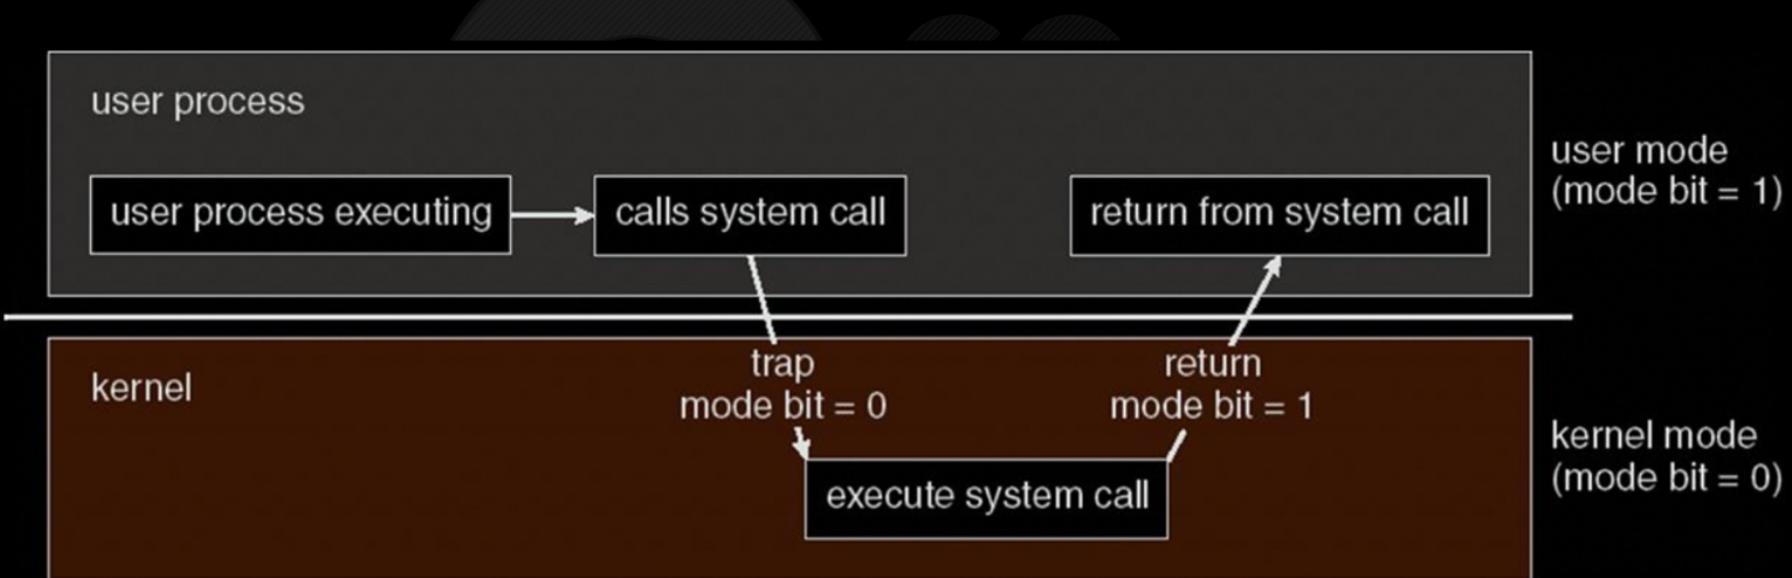
CLASSES



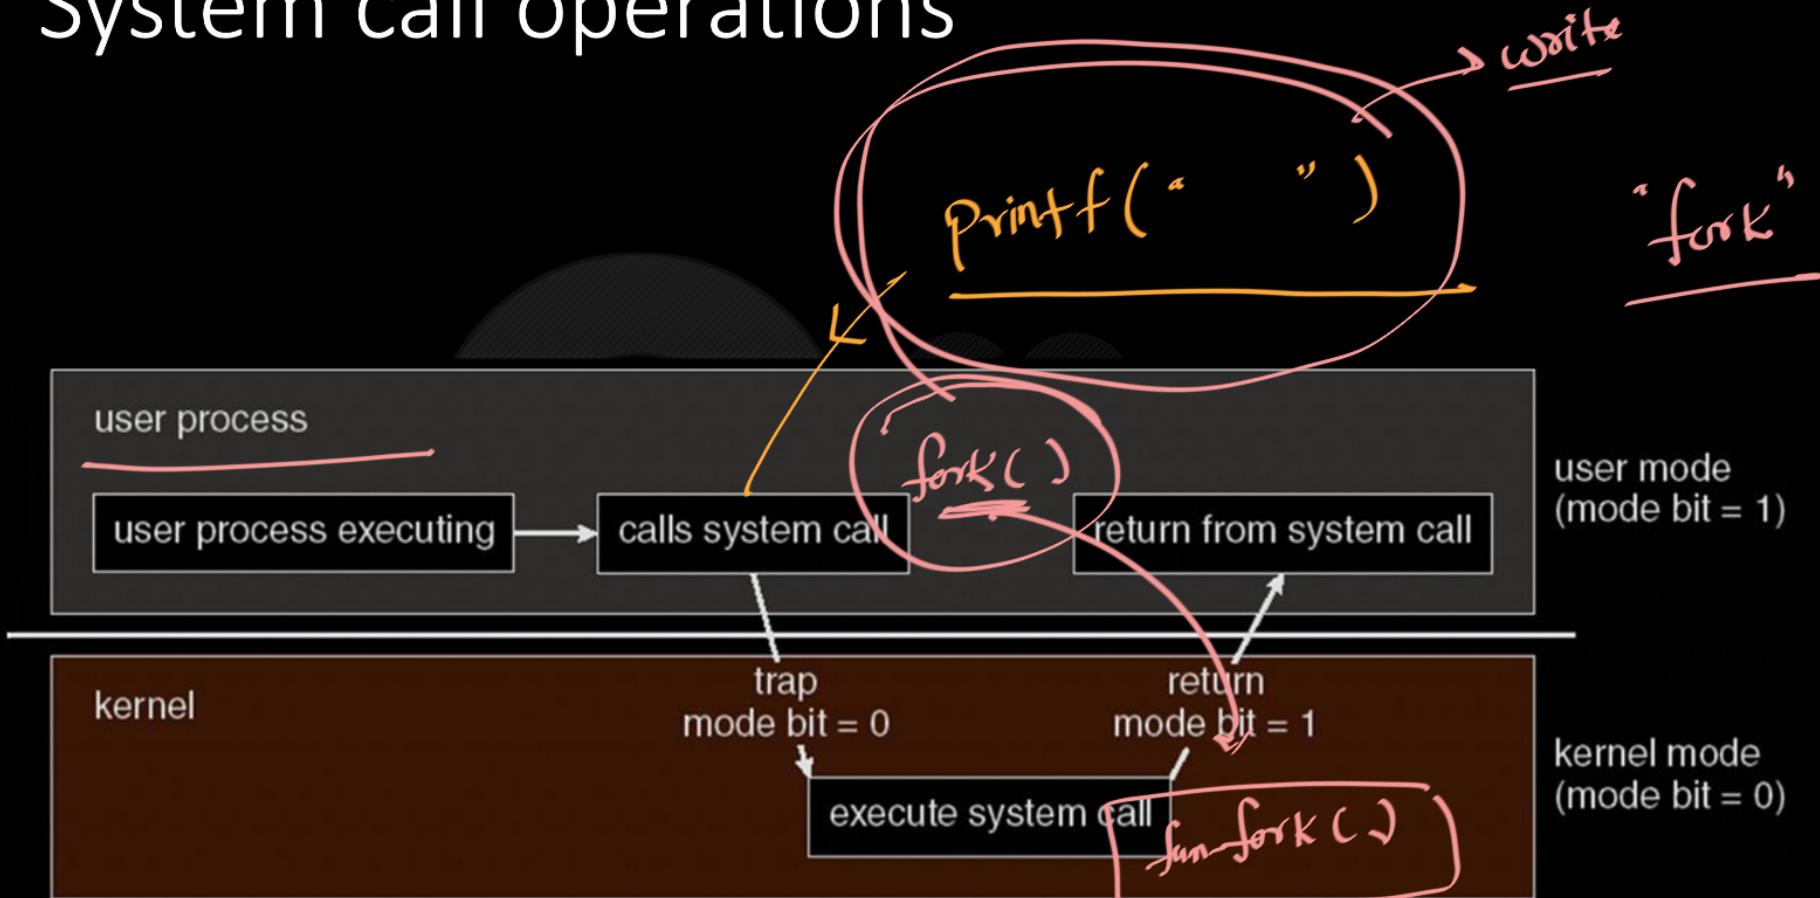
## system calls

- Note that we have blocked user (user program) from doing certain tasks to protect the illegal use of the system.
- Instead, user can now request OS to perform these tasks. These tasks will run privileged instructions → in OS code
- User program can initiate such request using a mechanism called system calls.
- Note that we need to change the mode here

# System call operations



## System call operations



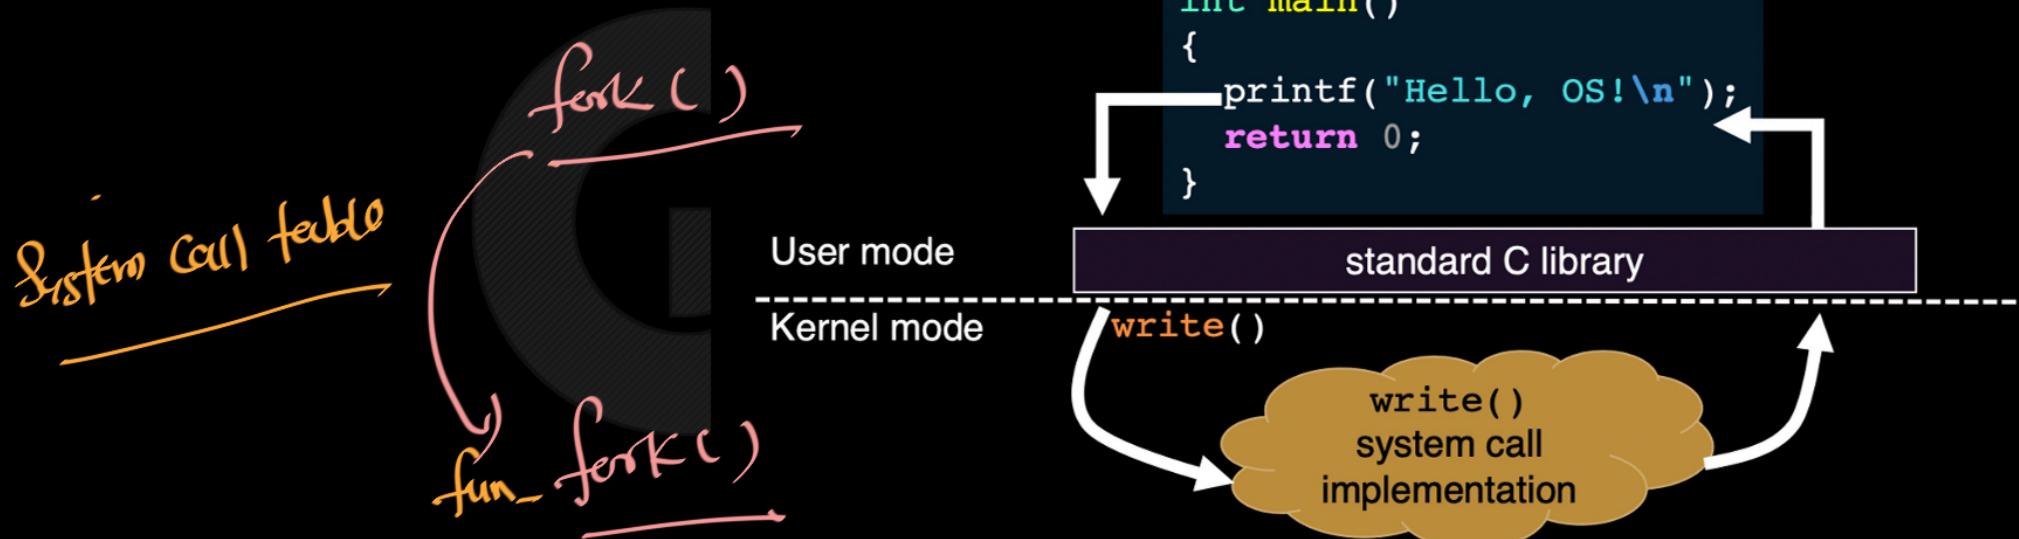


## Invoking System Calls

A user process cannot invoke system calls like regular functions because they are part of the kernel-space. Modern CPU provide a special instruction generally called a *software interrupt* , which can be performed by a user-space software and which causes the CPU to switch to the system mode and give the control to the kernel, just like when serving a hardware interrupt or an exception.

<https://courses.cs.washington.edu/courses/cse333/20su/lectures/lec-07/07-posix-20su.pdf>

## System call Interface





As a user can i ask directly from

monitor?

No  
==



GO  
CLASSES

A large, semi-transparent watermark of the Go Classes logo is centered on the page. It features a large, dark gray 'G' with a diagonal striped pattern, followed by the words "GO CLASSES" in a large, dark gray sans-serif font.

if i directly use monitor without asking then  
it is a problem

## Application Programming Interface

(API)

User processes usually do not directly invoke system calls. Instead they use something which is called an *Application Programming Interface* (an API, for short). It is a set of functions, variables and other software constructs that allow the processes to perform the most needed operations and also to cooperate with the operating system.

library

# Application Programming Interface

User processes usually do not directly invoke system calls. Instead they use something which is called an *Application Programming Interface* (an API, for short). It is a set of functions, variables and other software constructs that allow the processes to perform the most needed operations and also to cooperate with the operating system.

Linux has an API defined by the POSIX and SUS standards, which is inherited from the Unix system. Some of the functions from this library do not use any system calls (like the `strcpy()` function),

some of them are simple wrappings for a system call (like the `write()` function) and some of them perform other operations beside invoking a system call (like the `printf()` function).

Optional  
read

printf()  
↓  
System call  
Do some other this



# Operating Systems

4. System calls must be used to:
- (a) modify global variables
  - (b) call a user-written function
  - (c) write to a file
  - (d) All of the above

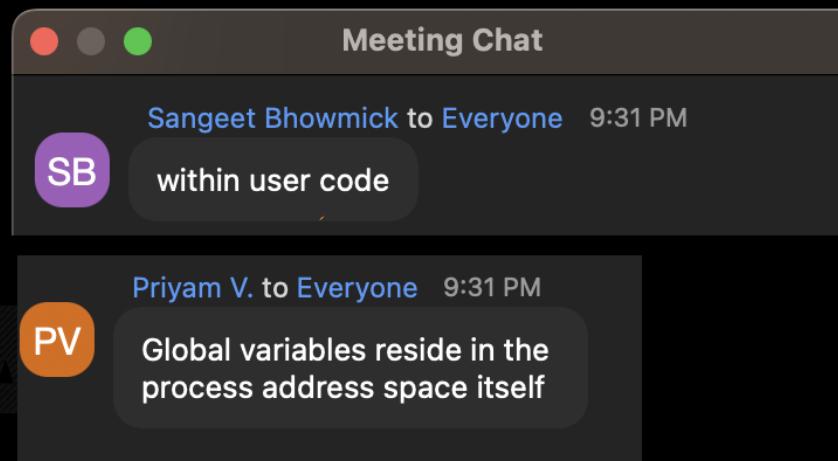
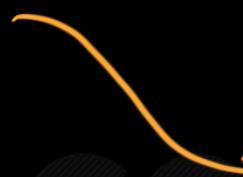
GO  
CLASSES



# Operating Systems

4. System calls must be used to:

- (a) modify global variables
- (b) call a user-written function
- (c) write to a file
- (d) All of the above



Answer C



System Call Execution



# System Call Execution

(or System call implementation)



(interrupt execution)  
in coA