



## Lecture 11

GO  
CLASSES



## Recap Question

For each of the following thread state transitions, say whether the transition is legal *and* how the transition occurs or why it cannot.

- i) Change from thread state BLOCKED to thread state RUNNING
- ii) Change from thread state RUNNING to thread state BLOCKED
- iii) Change from thread state RUNNABLE to thread state BLOCKED  
*(Ready)*



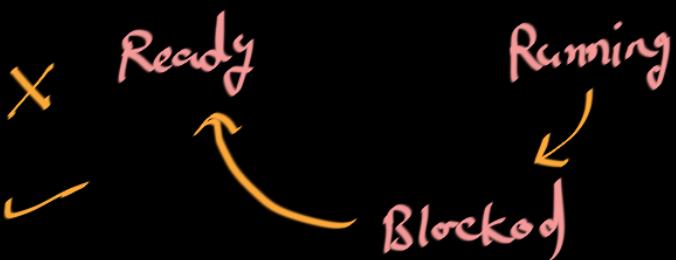
<https://www.cs.cornell.edu/courses/cs4410/2008fa/oldexams/midterm-solutions.pdf>



# Recap Question

For each of the following thread state transitions, say whether the transition is legal *and* how the transition occurs or why it cannot.

i) Change from thread state BLOCKED to thread state RUNNING



ii) Change from thread state RUNNING to thread state BLOCKED

iii) Change from thread state RUNNABLE to thread state BLOCKED

(Ready)

<https://www.cs.cornell.edu/courses/cs4410/2008fa/oldexams/midterm-solutions.pdf>



## Answer

For each of the following thread state transitions, say whether the transition is legal *and* how the transition occurs or why it cannot.

- i) Change from thread state BLOCKED to thread state RUNNING

**Illegal.** *The scheduler selects threads to run from the list of ready (or runnable) threads. A blocked thread must first be placed in the ready queue before it can be selected to run.*

- ii) Change from thread state RUNNING to thread state BLOCKED

**Legal.** *A running thread can become blocked when it requests a resource that is not immediately available (disk I/O, lock, etc).*

- iii) Change from thread state RUNNABLE to thread state BLOCKED

**Illegal.** *A thread can only transition to BLOCKED from RUNNING. It cannot execute any statements when still in a queue (i.e. the ready queue).*

SES



## First-Come First-Serve (FCFS)

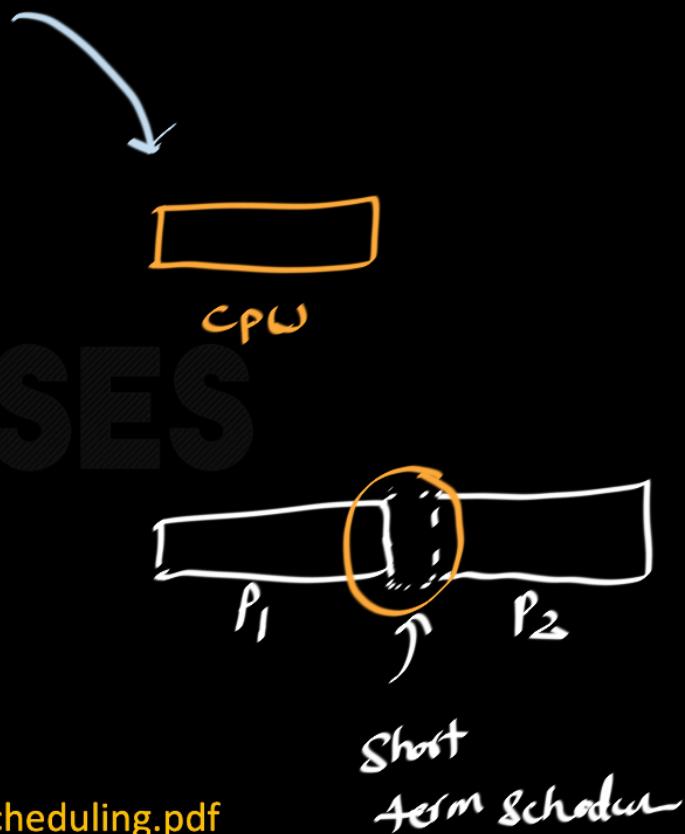
- By far the simplest CPU-scheduling algorithm

Process	Burst Time
$P_1$	24
$P_2$	3
$P_3$	3

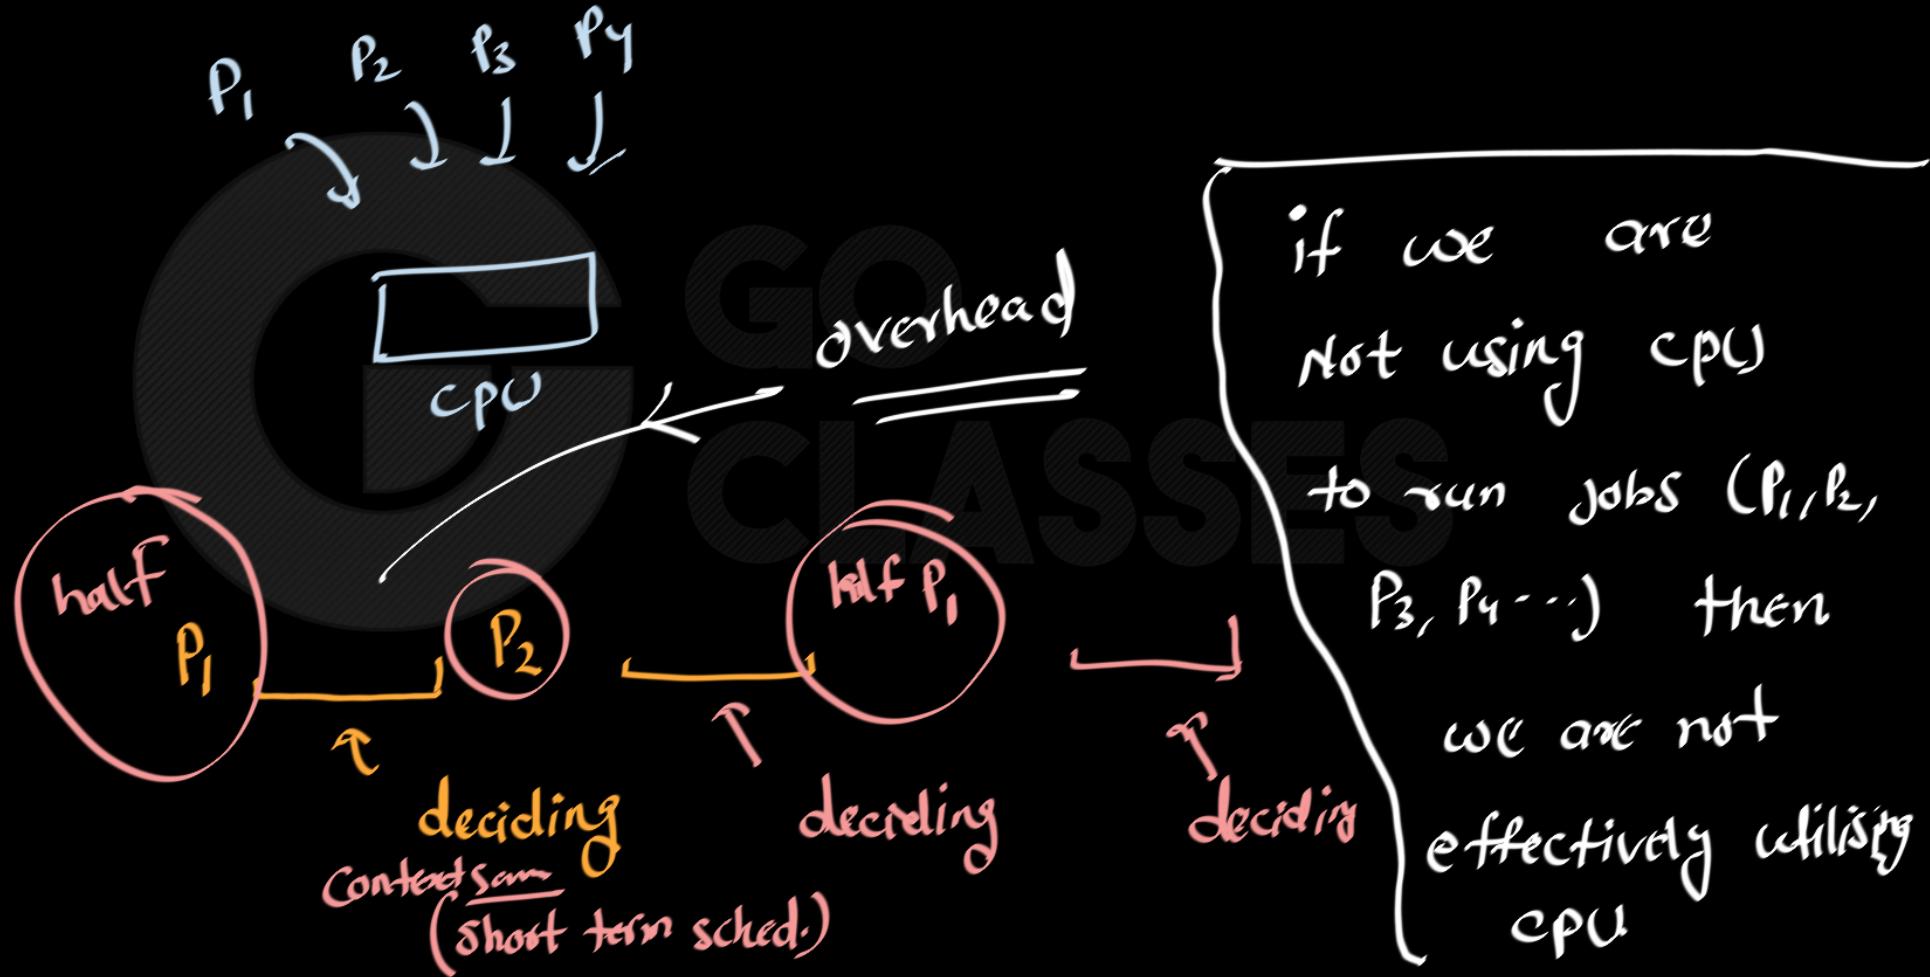


# FCFS vs Objective Functions

- Maximize CPU Utilization: Excellent
- Maximize Throughput: Highly dependent on first submitted job(s) duration
- Minimize Turnaround Time: Highly dependent on first submitted job(s) duration
- Minimize Waiting Time: Highly dependent on first submitted job(s) duration
- Minimize Response Time: Highly dependent on first submitted job(s) duration



[http://courses.ics.hawaii.edu/ReviewICS332/morea/060\\_Scheduling/ics332\\_scheduling.pdf](http://courses.ics.hawaii.edu/ReviewICS332/morea/060_Scheduling/ics332_scheduling.pdf)

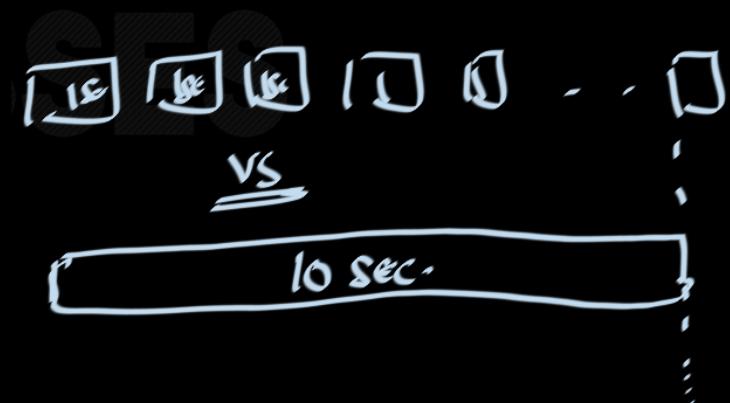




## FCFS vs Objective Functions

more clear in RR scheduling

- Maximize CPU Utilization: Excellent
- Maximize Throughput: Highly dependent on first submitted job(s) duration
- Minimize Turnaround Time: Highly dependent on first submitted job(s) duration
- Minimize Waiting Time: Highly dependent on first submitted job(s) duration
- Minimize Response Time: Highly dependent on first submitted job(s) duration



[http://courses.ics.hawaii.edu/ReviewICS332/morea/060\\_Scheduling/ics332\\_scheduling.pdf](http://courses.ics.hawaii.edu/ReviewICS332/morea/060_Scheduling/ics332_scheduling.pdf)



in 10 sec.

=

we are

finishing

10 jobs

throughput =

$$10 \text{ jobs} / 10 \text{ sec}$$

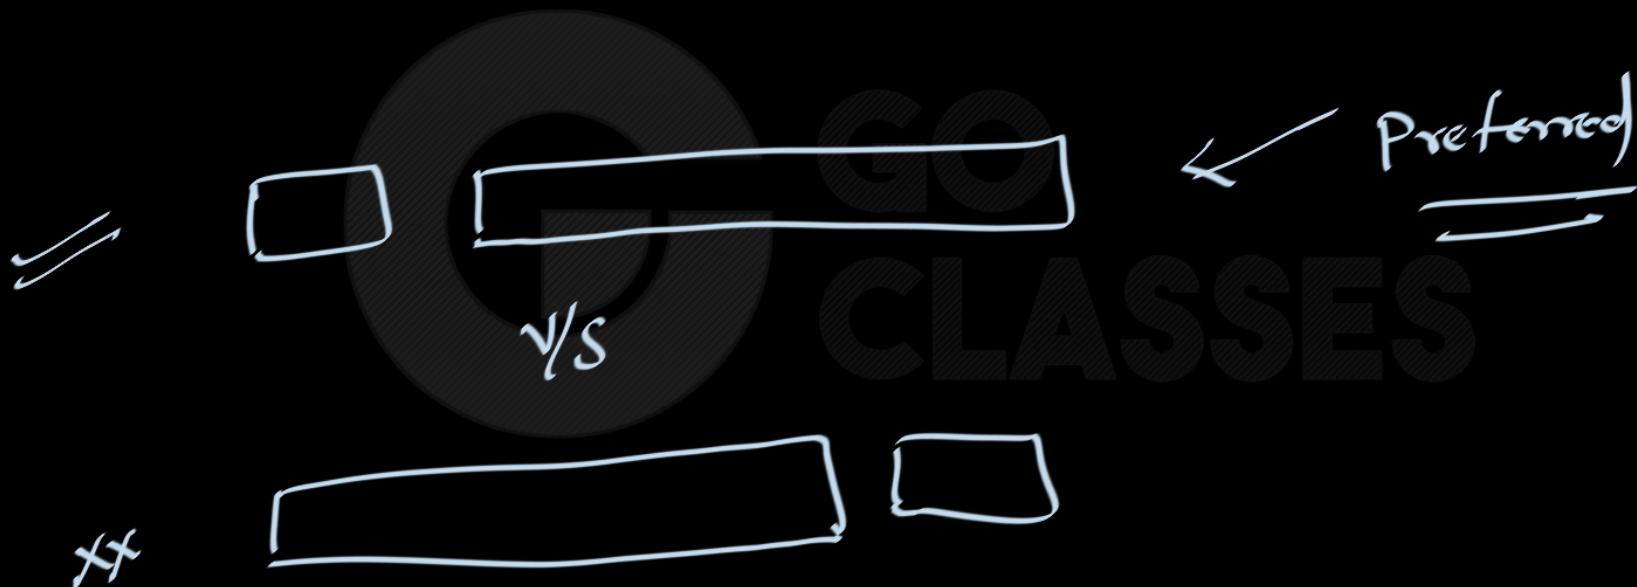
$$1 \text{ job/sec}$$

10 sec.

$$\text{throughput} = \frac{1 \text{ job}}{10 \text{ sec}} = \frac{1}{10} \text{ job}$$



## Shortest Job First (SJF)





## Shortest-Job-First Scheduling (SJF)

- When the CPU is available, it is assigned to the process that has the smallest next CPU burst.
- If the next CPU bursts of two processes are the same, FCFS scheduling is used to break the tie.



# Operating Systems

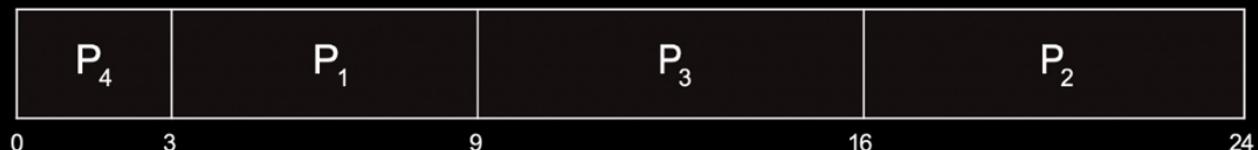
<u>Process</u>	<u>Burst Time</u>
$P_1$	6
$P_2$	8
$P_3$	7
$P_4$	3



# Operating Systems

<u>Process</u>	<u>Burst Time</u>
$P_1$	6
$P_2$	8
$P_3$	7
$P_4$	3

- SJF scheduling chart



ES

- Average waiting time =  $(3 + 16 + 9 + 0) / 4 = 7$

## Non-preemptive SJF: Example

Process	Duration	Arrival Time
P1	6	0
P2	8	0
P3	7	0
P4	3	0

Do it yourself



P4 waiting time: 0  
P1 waiting time: 3  
P3 waiting time: 9  
P2 waiting time: 16

The total time is: 24  
The average waiting time (AWT):  
 $(0+3+9+16)/4 = 7$

SES



## Comparing to FCFS

Process	Duration	Arrival Time
P1	6	0
P2	8	0
P3	7	0
P4	3	0

Do it yourself



P1 waiting time: 0

P2 waiting time: 6

P3 waiting time: 14

P4 waiting time: 21

The total time is the same.

The average waiting time (AWT):

$$(0+6+14+21)/4 = 10.25$$

(compared to 7)



# Predicting length of the next CPU burst

Let  $T_n$  be the length of the  $n^{th}$  CPU burst, and let  $S_{n+1}$  be our predicted value for the next CPU burst.

We define,

$$S_{n+1} = \alpha (T_n) + (1 - \alpha) * S_n$$

The value of  $T_n$  contains our most recent information, while  $S_n$  stores the past history.

- if  $\alpha = 0$ , then  $S_{n+1} = S_n$ , and recent history has no effect.
- if  $\alpha = 1$ , then  $S_{n+1} = T_n$ , and only the most recent CPU burst matters (history is assumed to be old and irrelevant).
- More commonly,  $\alpha = 1/2$ , so recent history and past history are equally weighted.



# Operating Systems

To understand the behavior of the exponential average, we can expand the formula for  $S_{n+1}$  by substituting for  $S_n$  to find

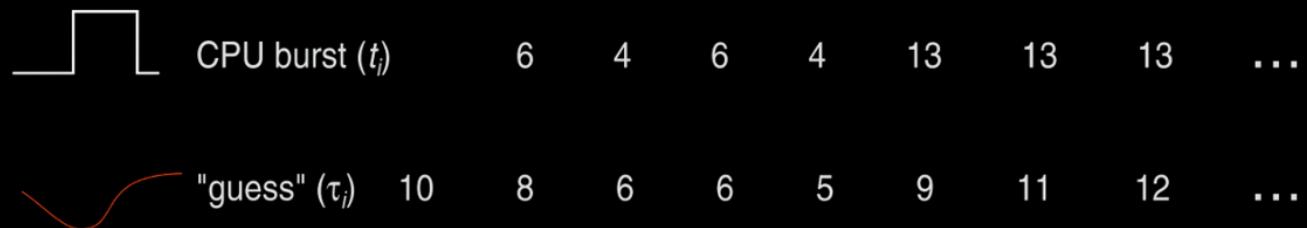
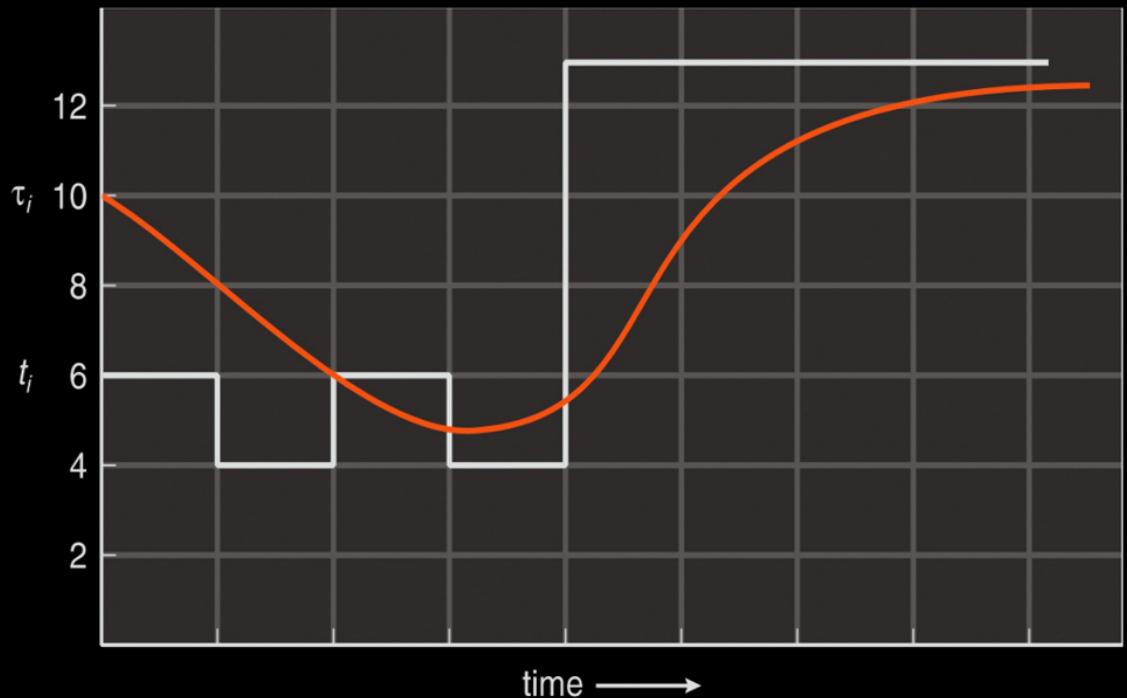
$$S_{n+1} = \alpha T_n + (1 - \alpha)\alpha T_{n-1} + \cdots + (1 - \alpha)^j \alpha T_{n-j} + \cdots + (1 - \alpha)^{n+1} S_0$$

Typically,  $\alpha$  is less than 1 . As a result,  $(1 - \alpha)$  is also less than 1 , and each successive term has less weight than its predecessor.





# Operating Systems





## Question

we initially had jobs with actual run times 8, 7, 4, and 16.

← Arrival time = 0

Assuming no new jobs entered the queue, we'd schedule them in increasing order of their times, namely 4, 7, 8, 16 (since that's what SJF means--shortest job first).

Calculate the exponential averaging  $S_4$  with  $S_1 = 10$ ,  $\alpha = 0.5$  and the algorithm is SJF.

- 1. 9
- 2. 8
- 3. 7.5
- 4. None

initial queue  
        

$$S_n = \frac{1}{2} S_{n-1} + \frac{1}{2} t_{n-1}$$

$$S_2 =$$



# Question

we initially had jobs with actual run times 8, 7, 4, and 16.

← Arrival time = 0

Assuming no new jobs entered the queue, we'd schedule them in increasing order of their times, namely 4, 7, 8, 16 (since that's what SJF means--shortest job first).

w n T real

Calculate the exponential averaging  $S_4$  with  $S_1 = 10$ ,  $\alpha = 0.5$  and the algorithm is SJF.

- 1. 9
- 2. 8
- 3. 7.5
- 4. None

initial queue

$$S_n = \frac{1}{2} S_{n-1} + \frac{1}{2} t_{n-1}$$

$$S_2 = \frac{1}{2} 10 + \frac{1}{2} 4 = 7$$

$$S_3 =$$



# Operating Systems

## Answer

$$S_2 = \frac{1}{2}(4 + 10) = 7 \quad \text{after the first job has run}$$

$$S_3 = \frac{1}{2}(7 + 7) = 7 \quad \text{after the second job has run}$$

$$S_4 = \frac{1}{2}(8 + 7) = \underline{\underline{7.5}} \quad \text{after the third job has run}$$

so the fourth job would be given a slice of 7.5, which is choice (3).



## Optional Read from Galvin

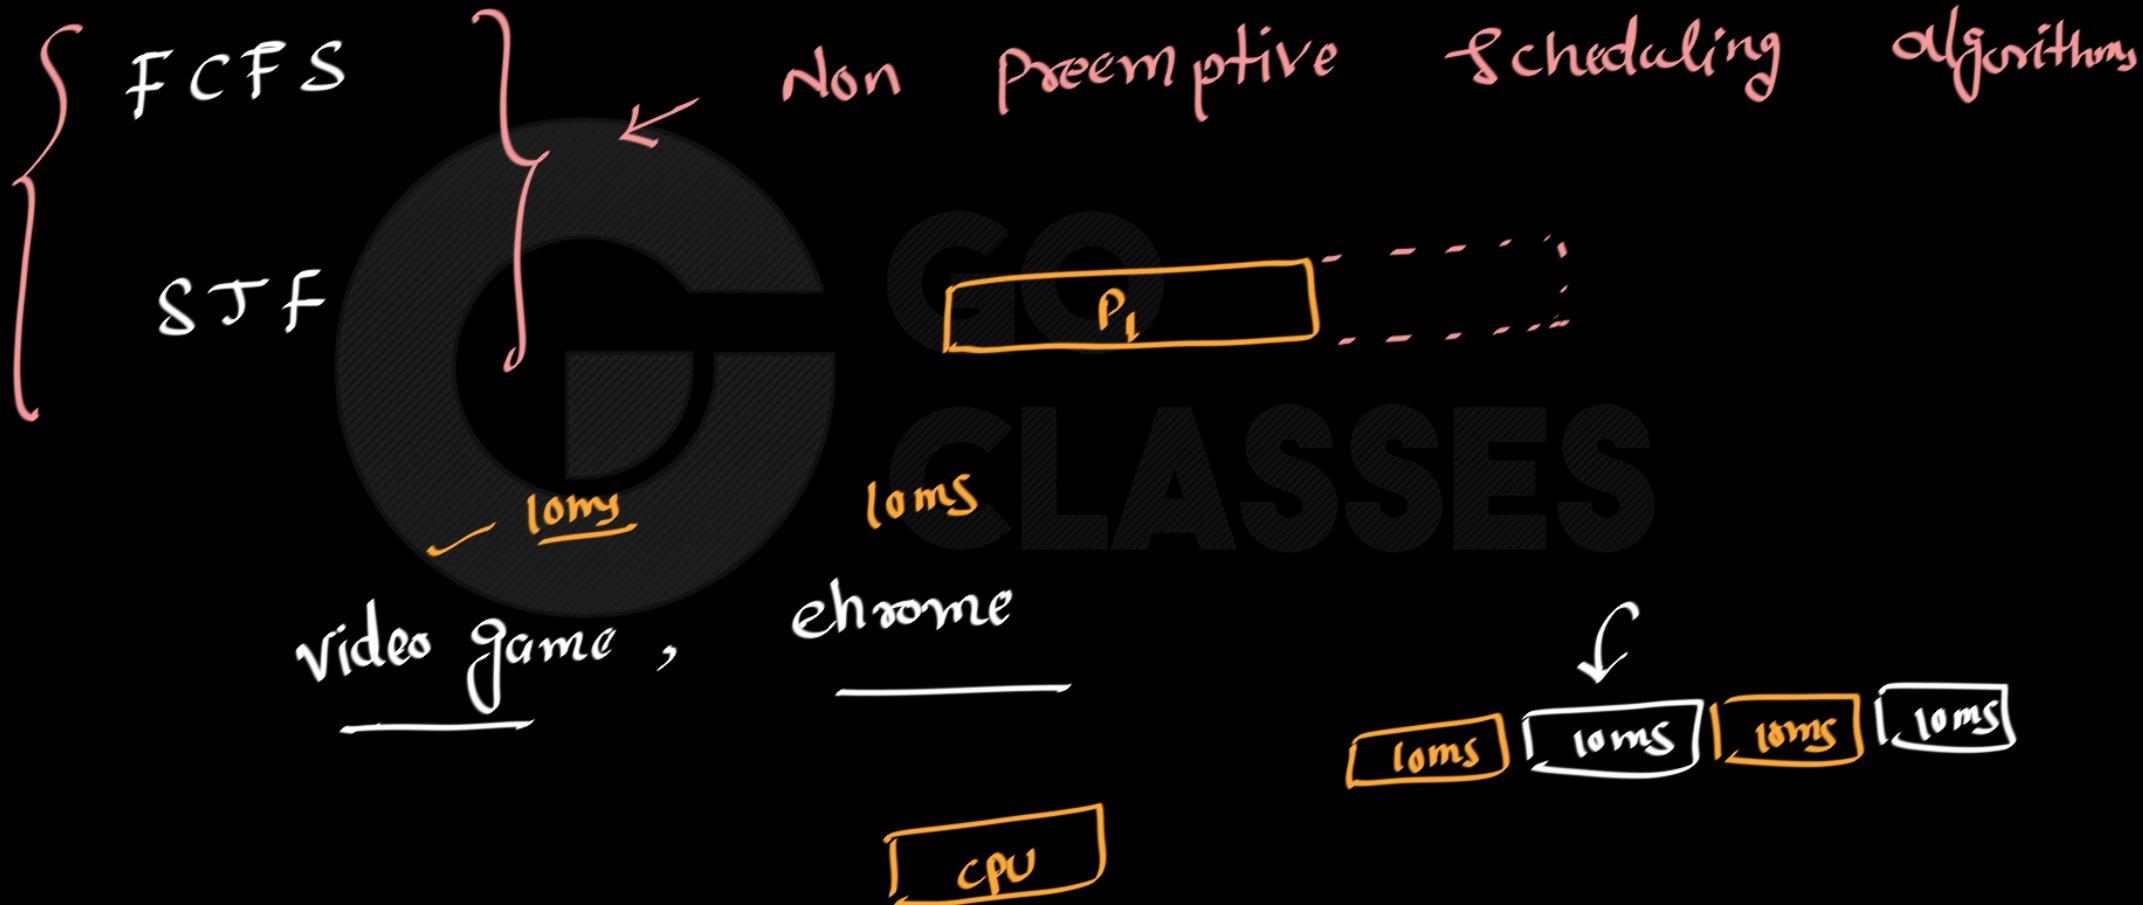
The SJF algorithm can be either preemptive or nonpreemptive. The choice arises when a new process arrives at the ready queue while a previous process is still executing. The next CPU burst of the newly arrived process may be shorter than what is left of the currently executing process. A preemptive SJF algorithm will preempt the currently executing process, whereas a nonpreemptive SJF algorithm will allow the currently running process to finish its CPU burst. Preemptive SJF scheduling is sometimes called **shortest-remaining-time-first** scheduling.



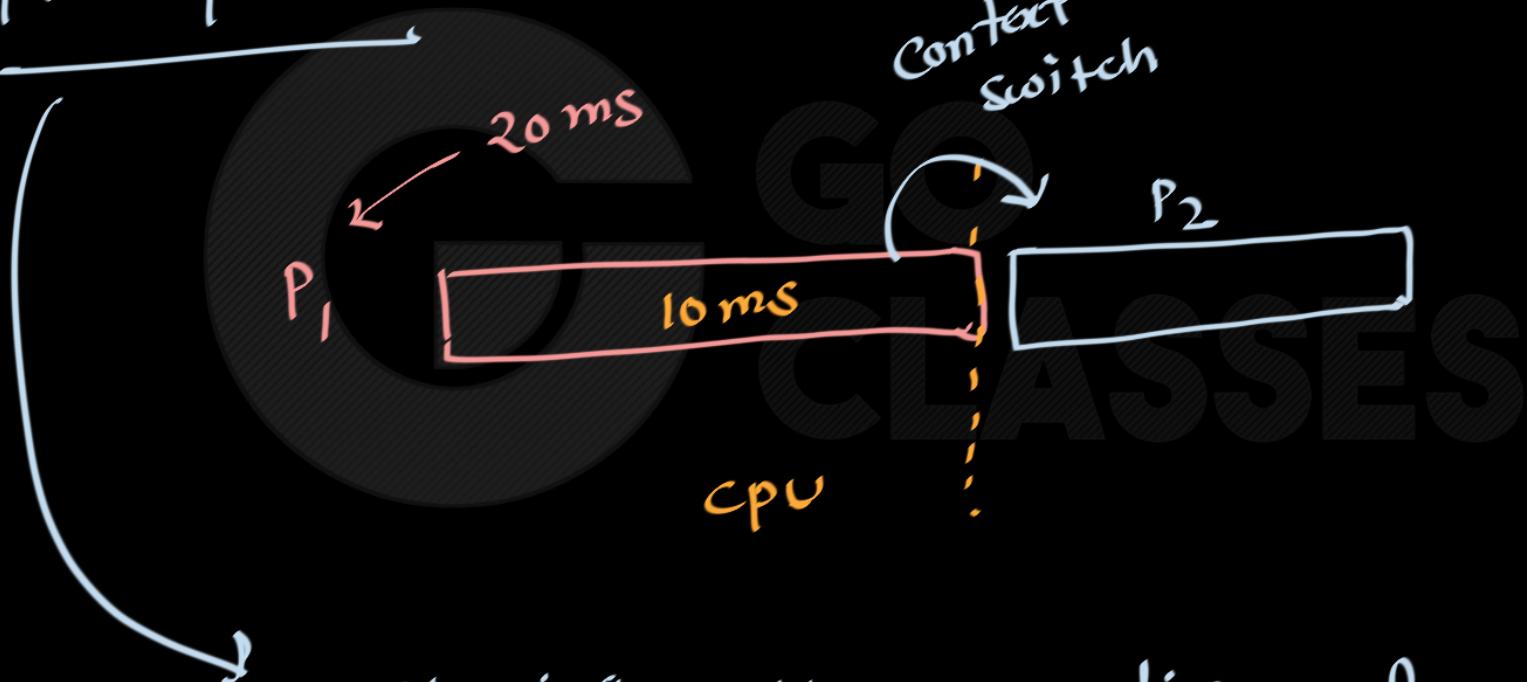
# Operating Systems



ROUND ROBIN      GO CLASSES



Preemption :



Stopping the execution of one process  
and running another process.



for preemption, we need to have context switch





if we keep on switching between video game

and chrome browser



it will provide an illusion to the user that

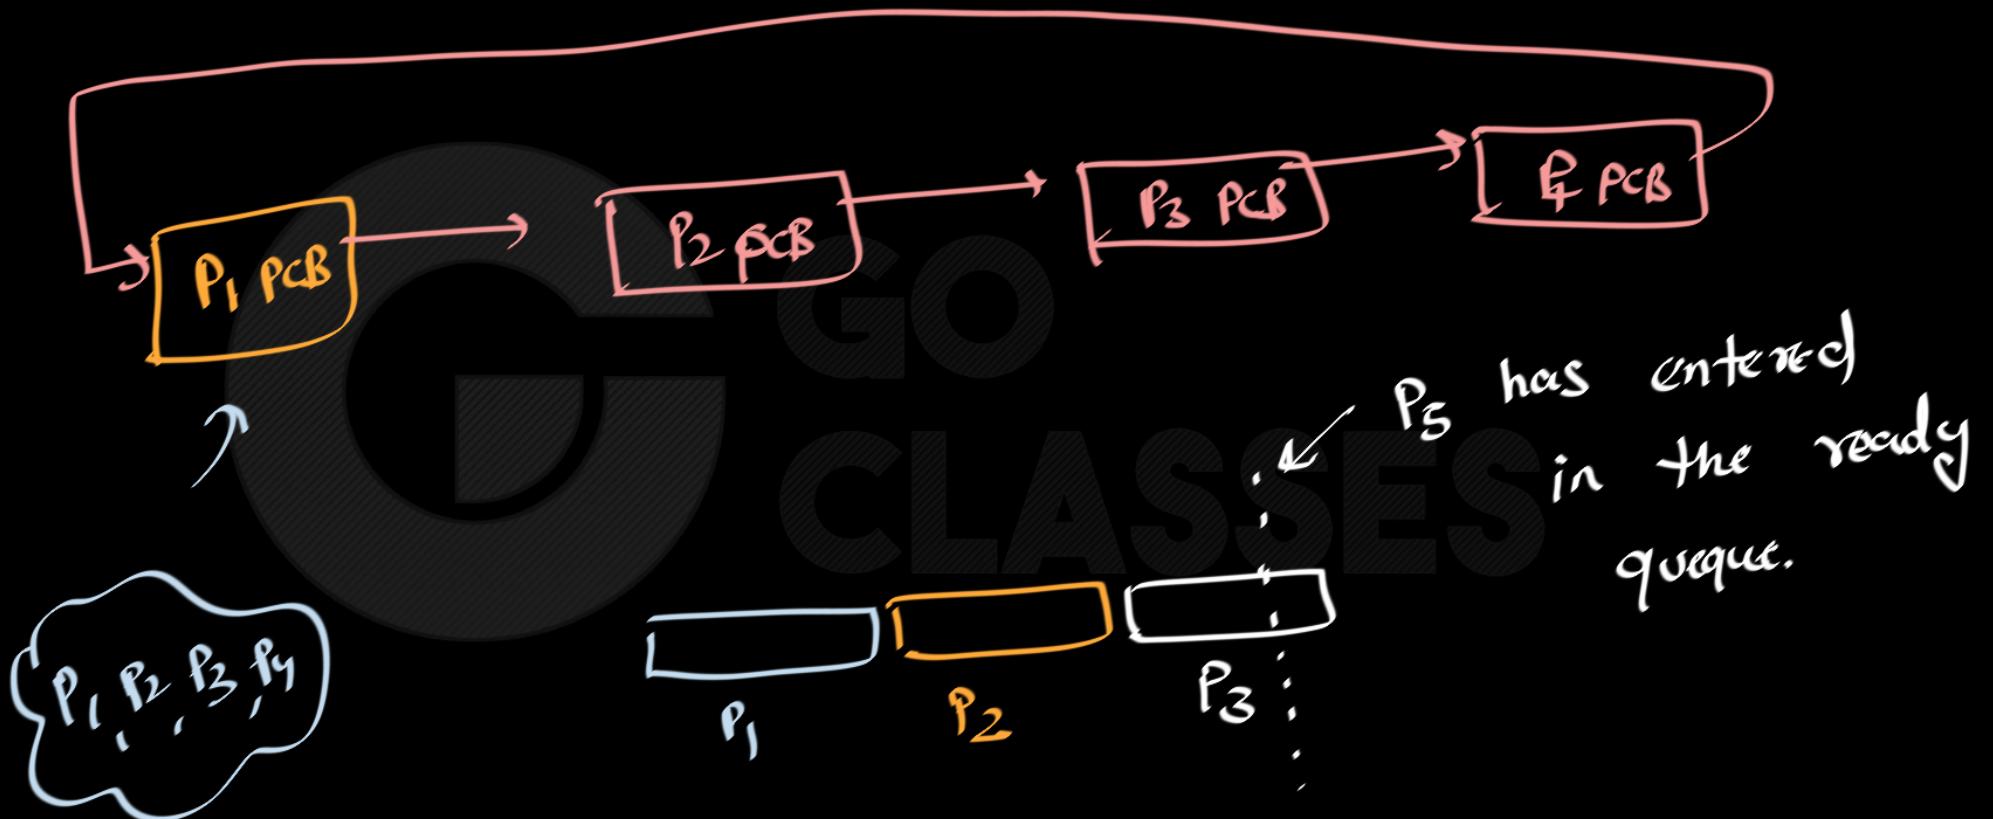
both are running simultaneously

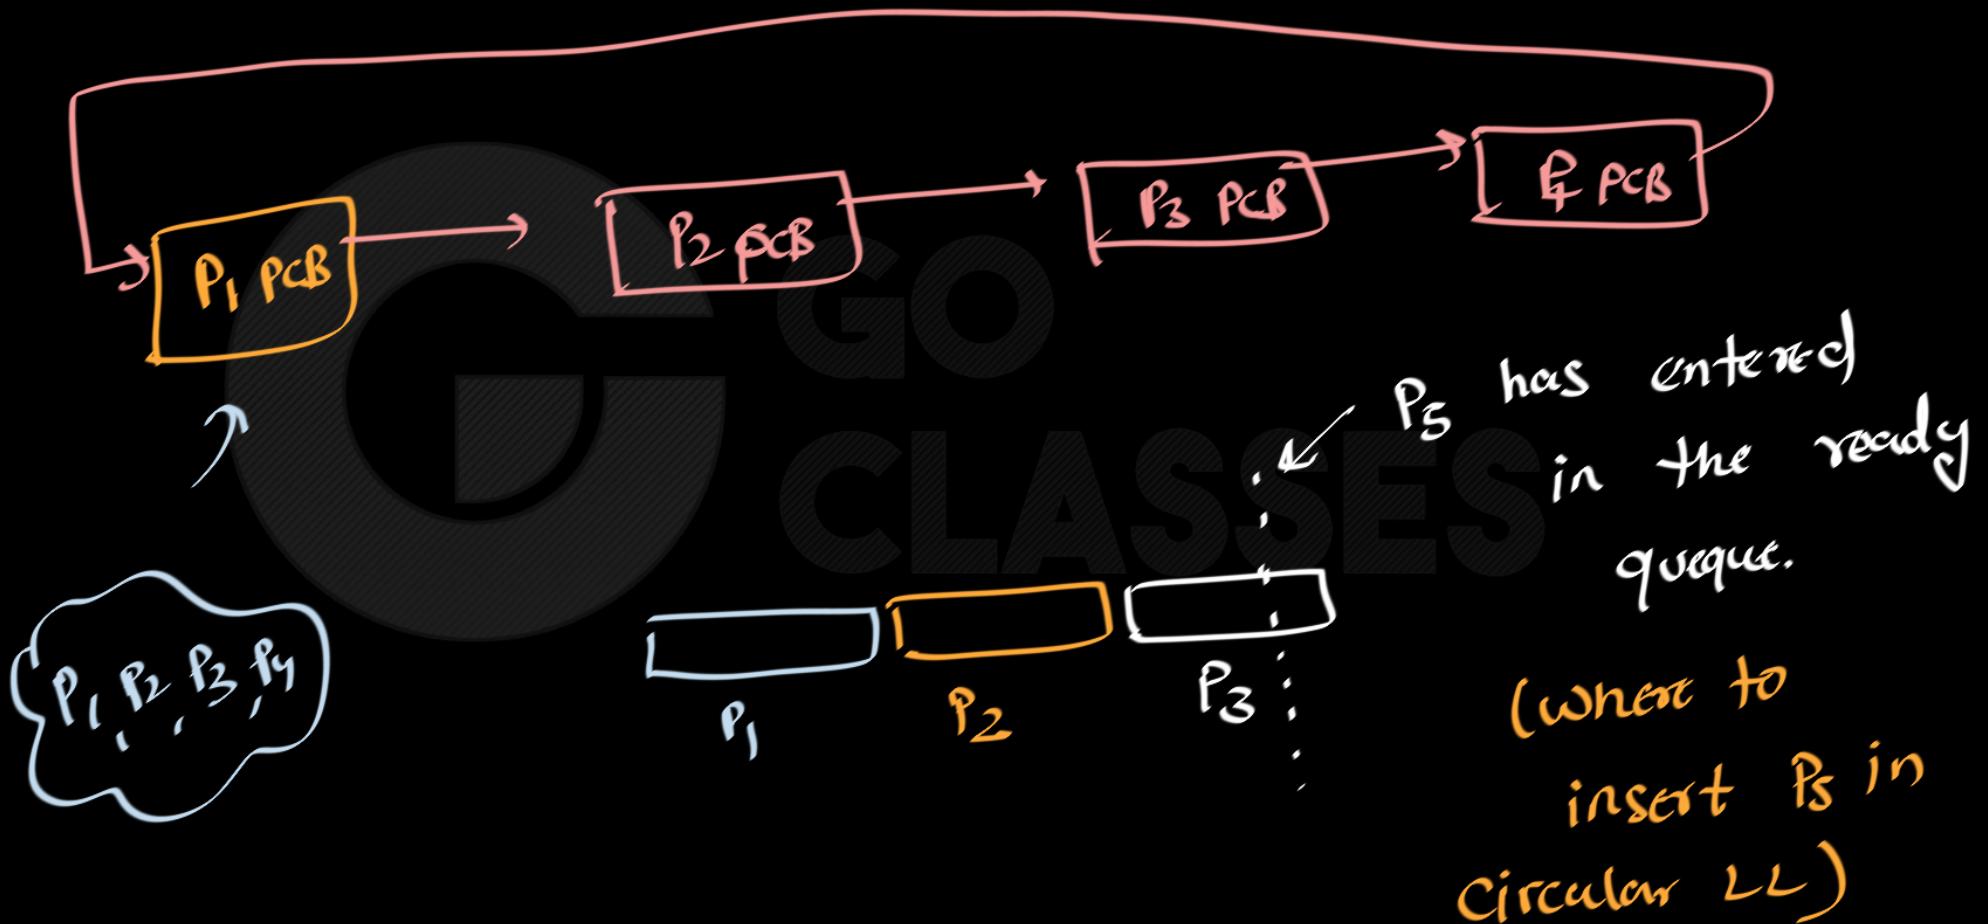
(in some cases it is desirable to switch between  
of processes)

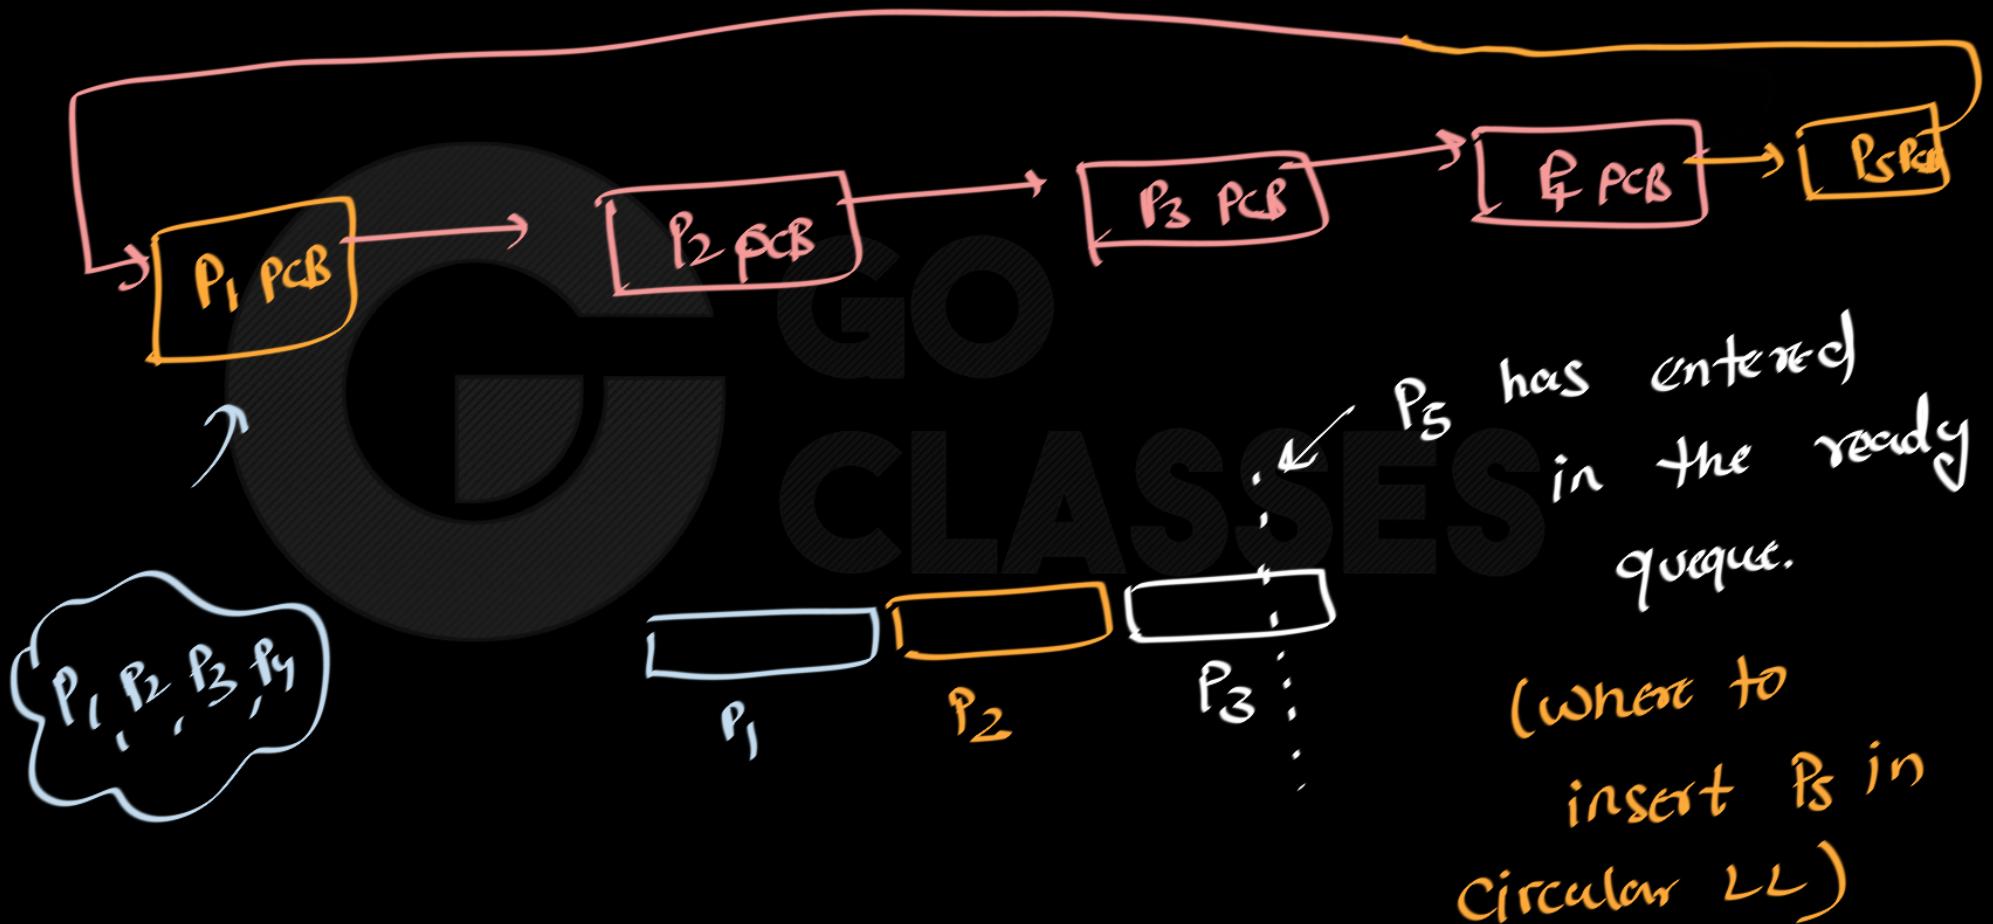
Ready queue

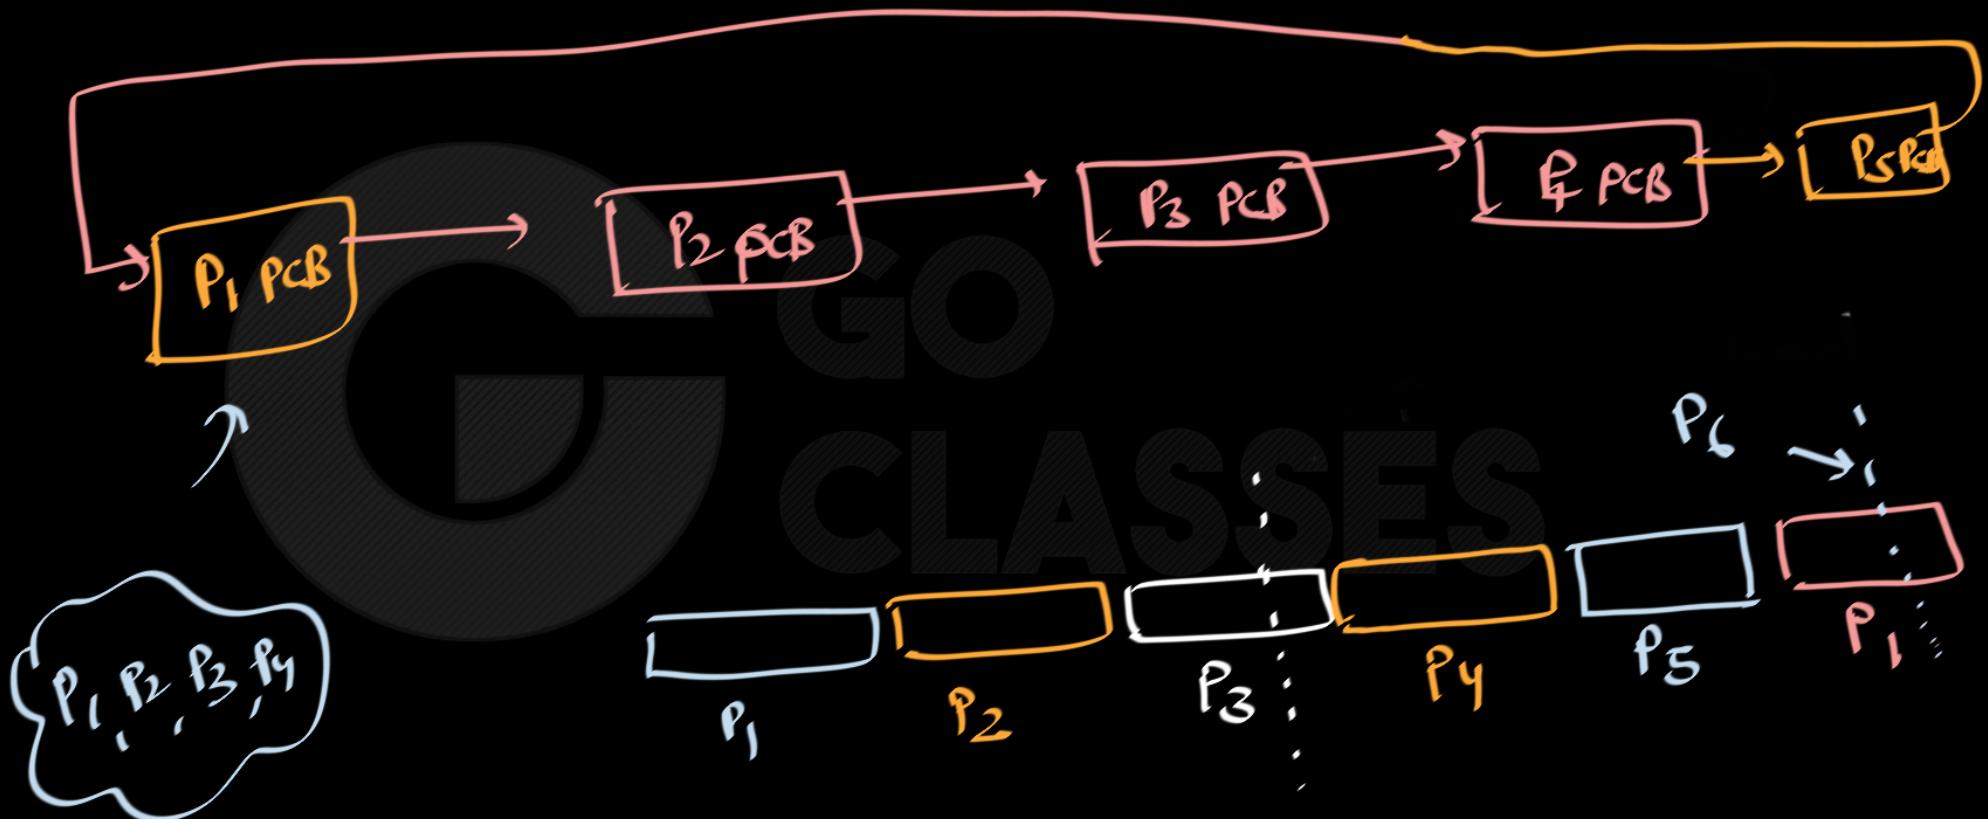


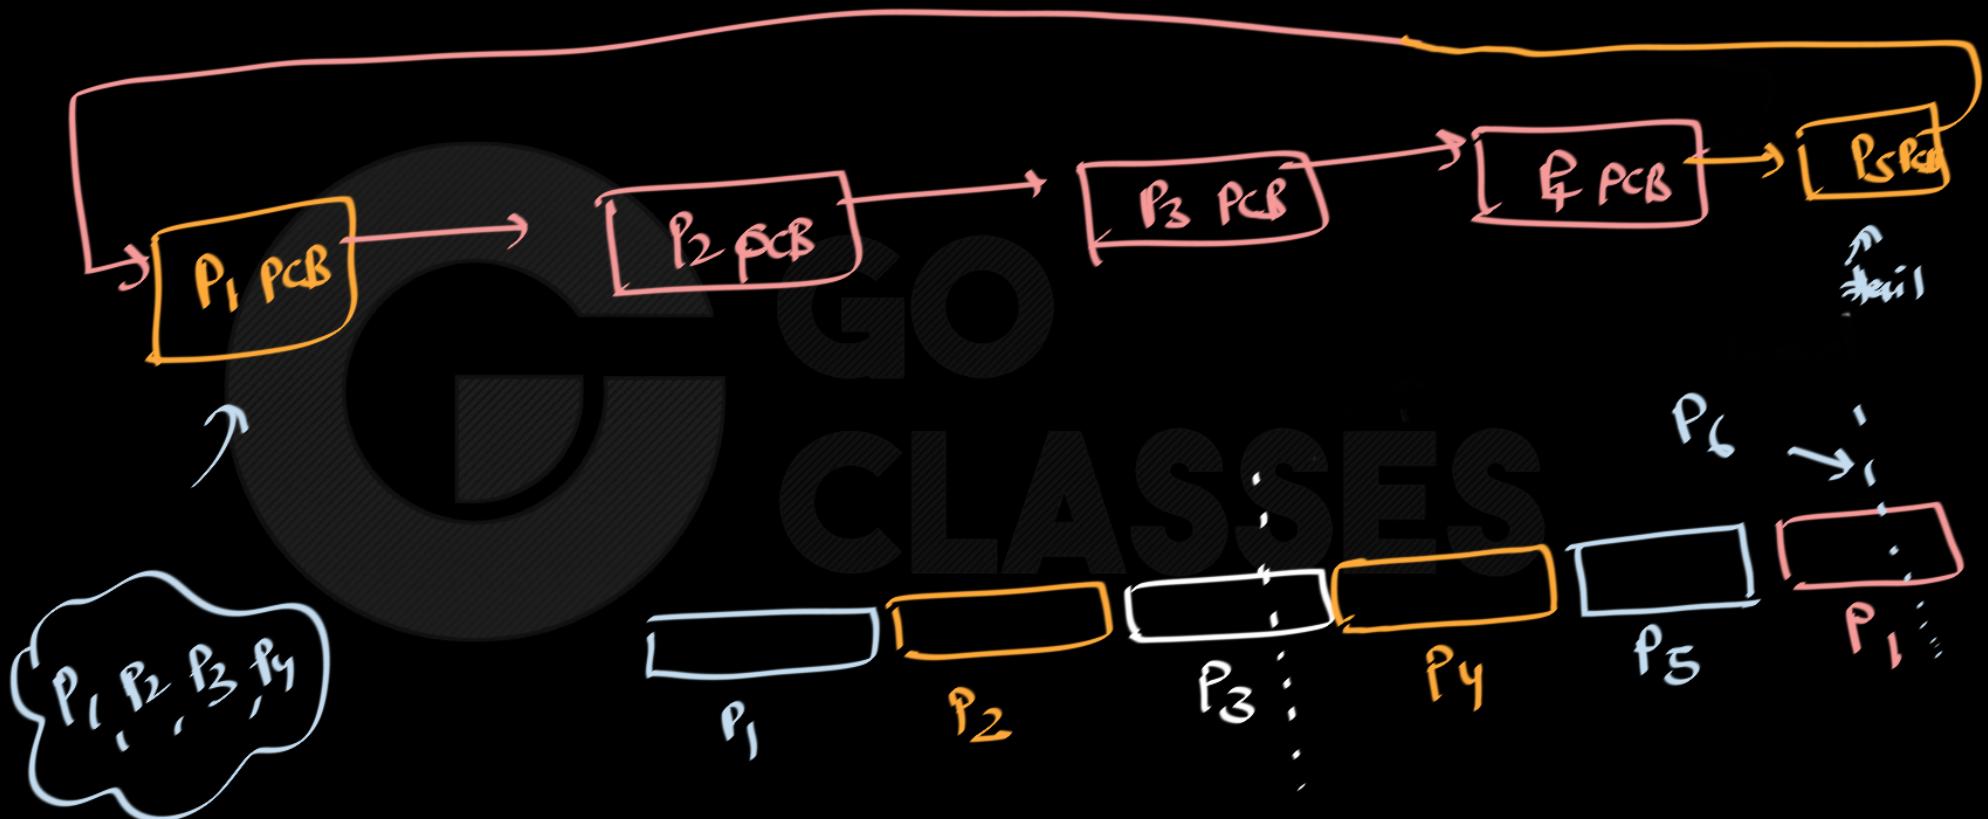
what data structure we can use for processes?

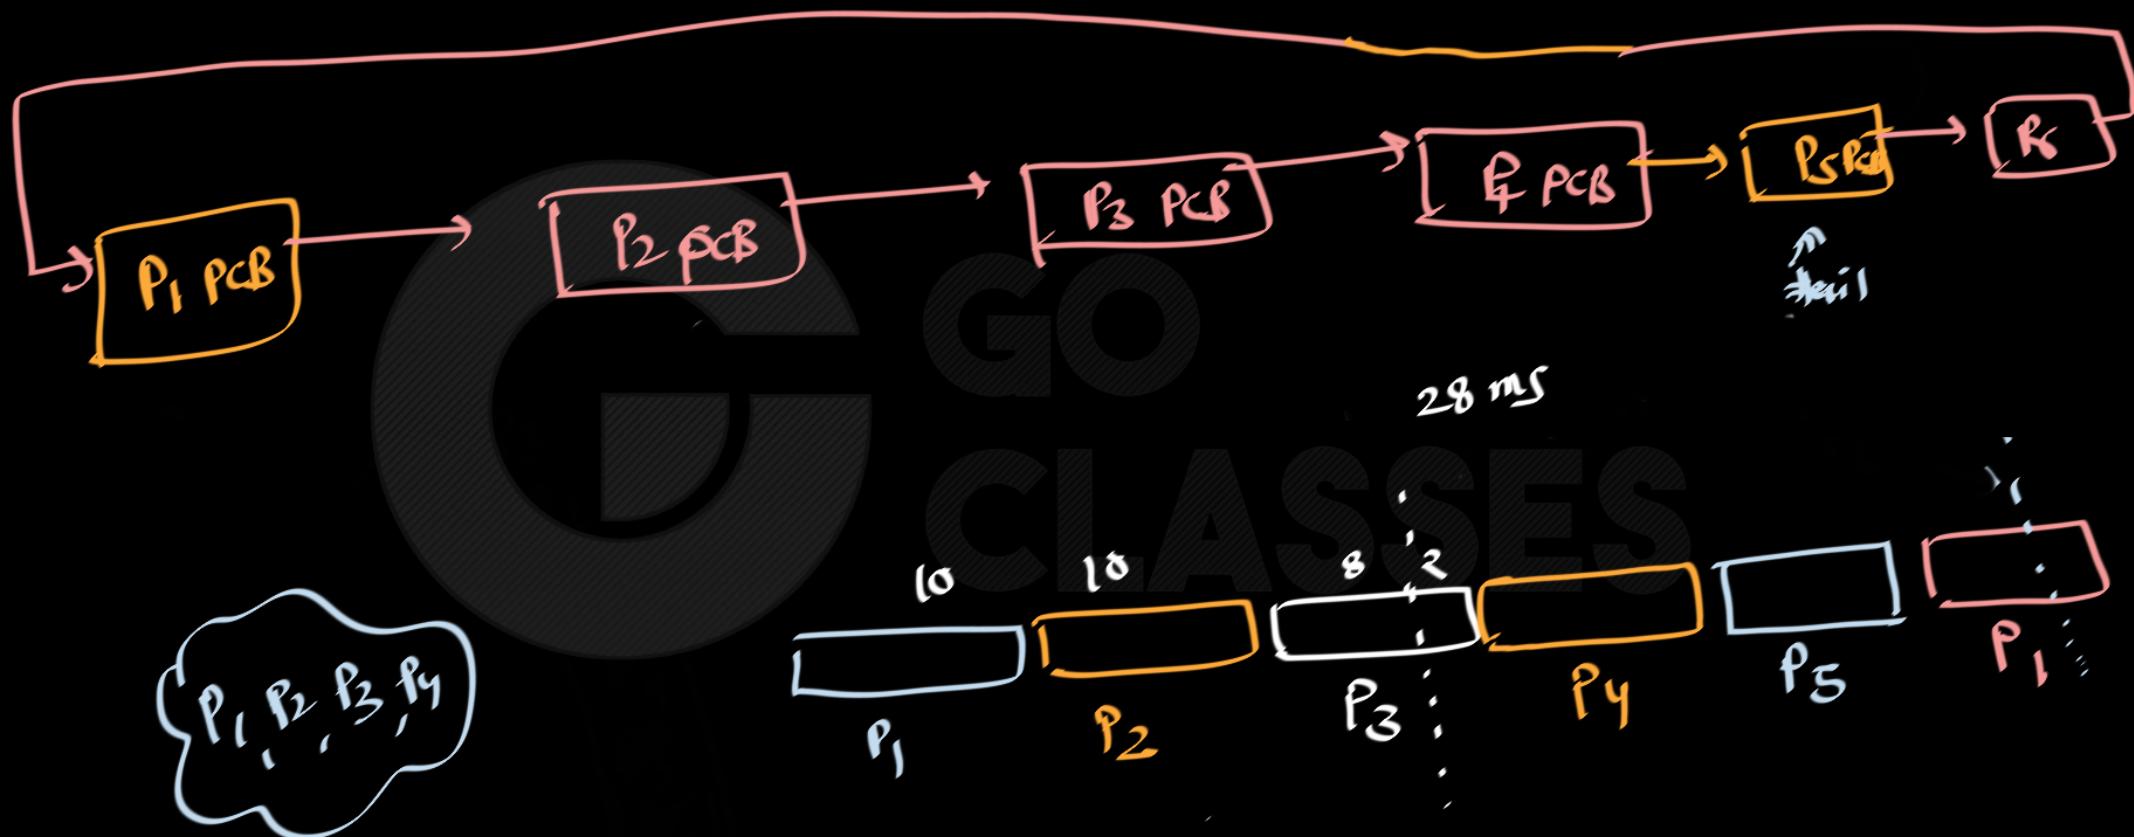










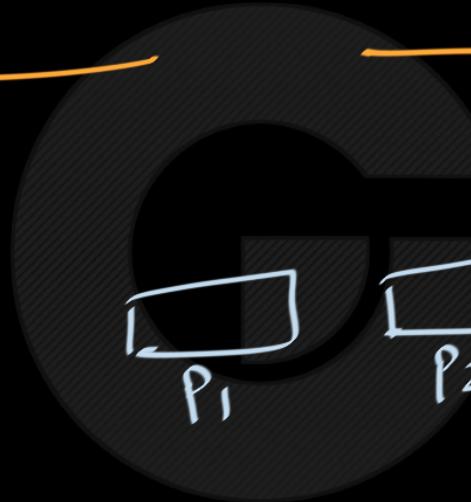


{ whatever initial processes are there in the ready queue make a circular list out of that

{ start execution of processes

{ if some process is coming in between the park it at the end of circular list.

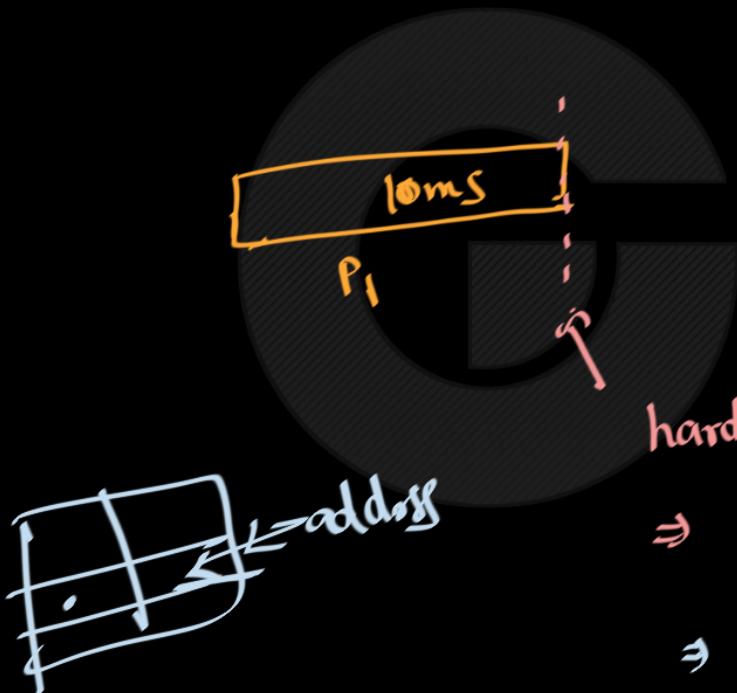
## Round Robin scheduling



follows

"quanta"  $\approx$  this timer is maintained  
 $\equiv$  by hardware.

time quantum = 10 ms



GO

CLASSES

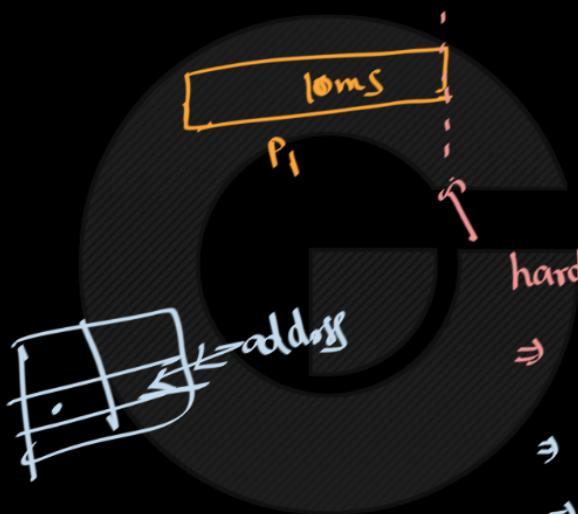
hardware timer will be triggered

⇒ this is one interrupt

⇒ Interrupt service routine will be checked

⇒ this in turn will lead to short term scheduler code

time quanta = 10 ms



hardware timer will be triggered

⇒ this is one interrupt

⇒ Interrupt service routine will be checked

⇒ this in turn will lead to short term  
scheduler code

→ generally very very fast because  
decision about new process to be scheduled  
is very easy.



# Operating Systems

Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

Arrival time = 0

$P_1$



# Operating Systems

Example of RR with Time Quantum = 4

Process

$P_1$

$P_2$

$P_3$

Burst Time

~~24~~ 20

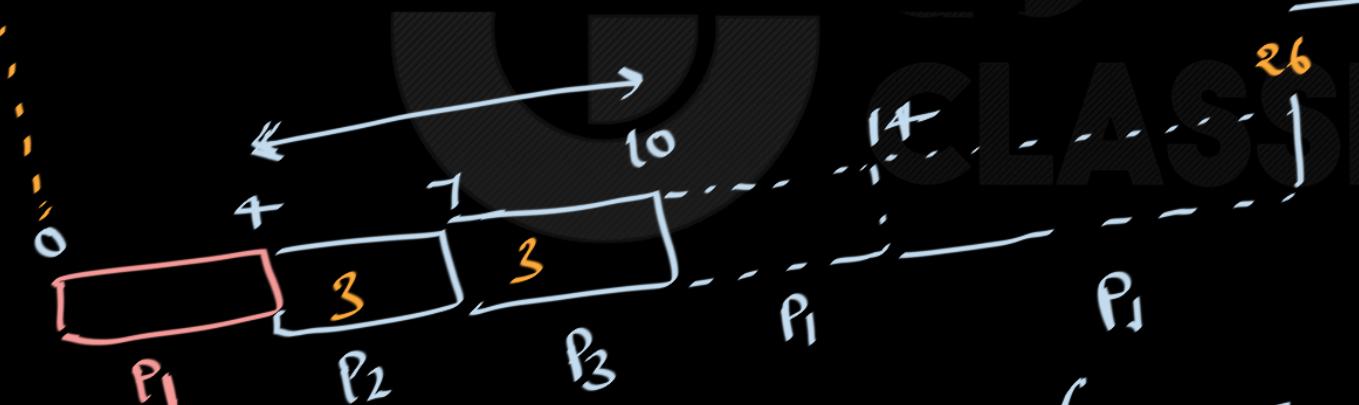
~~30~~ 10

~~30~~ 14

Arrival time = 0

$$\text{Avg. TAT} = \frac{26}{3}$$

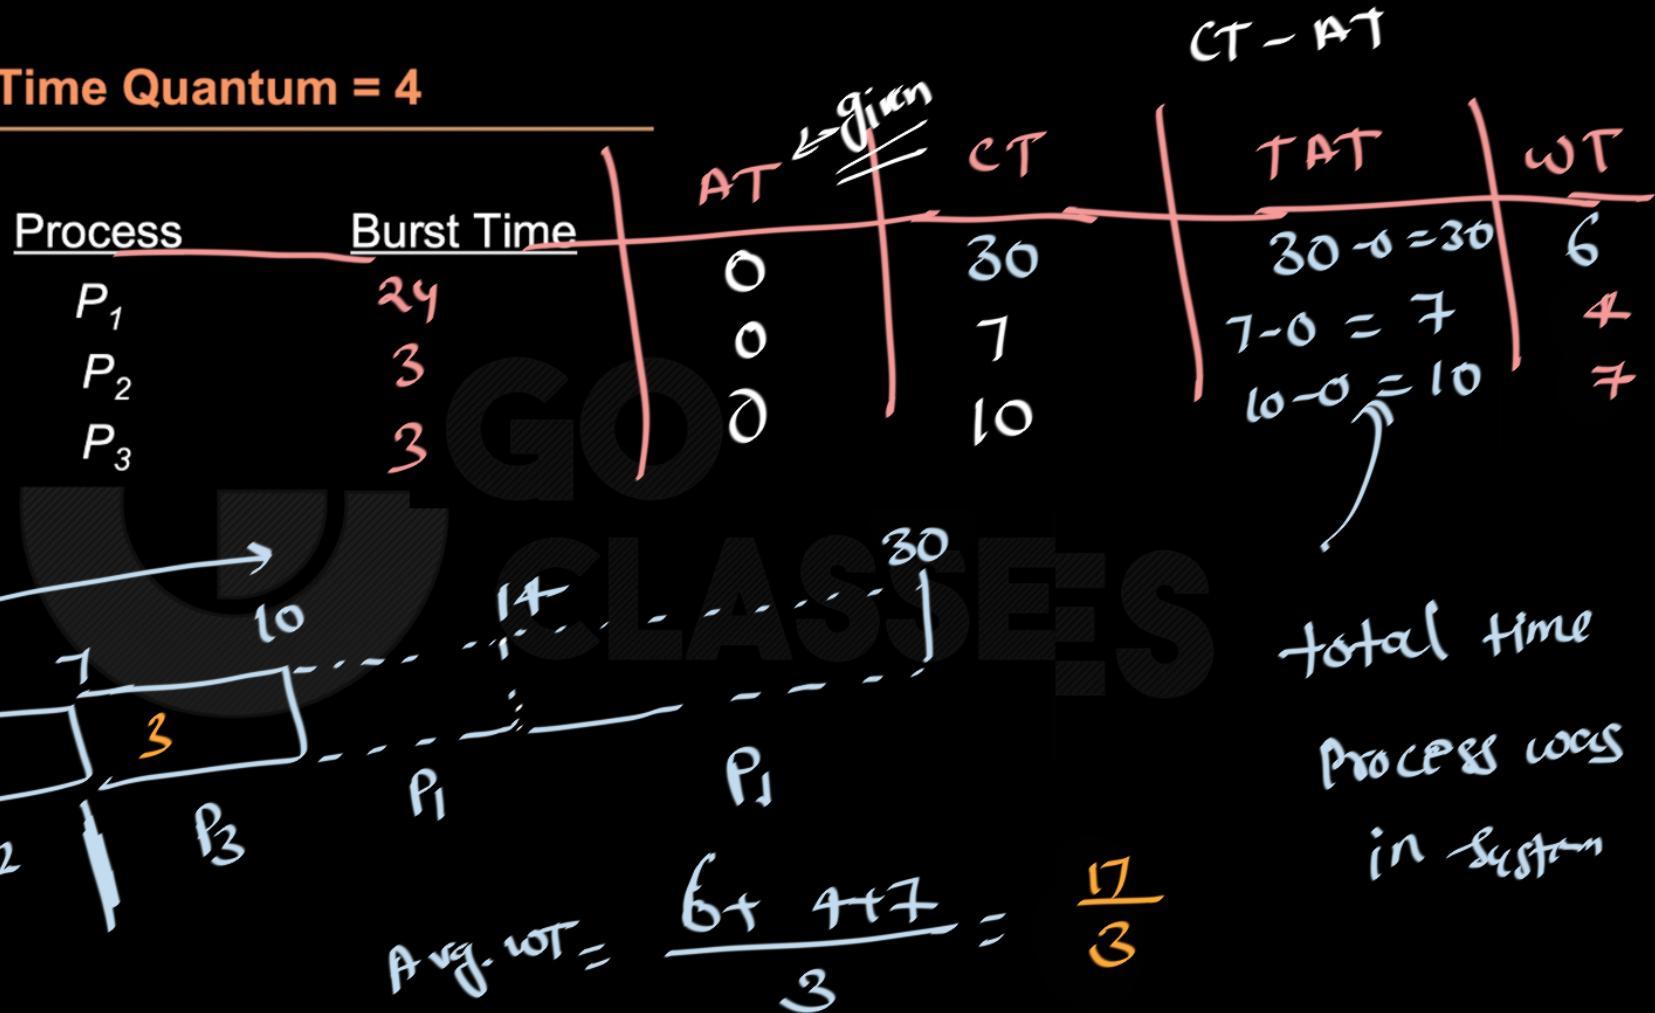
$$\frac{26}{3} = CT - AT$$



$$\text{Avg. WT} = \frac{6+4+7}{3} = \frac{17}{3}$$

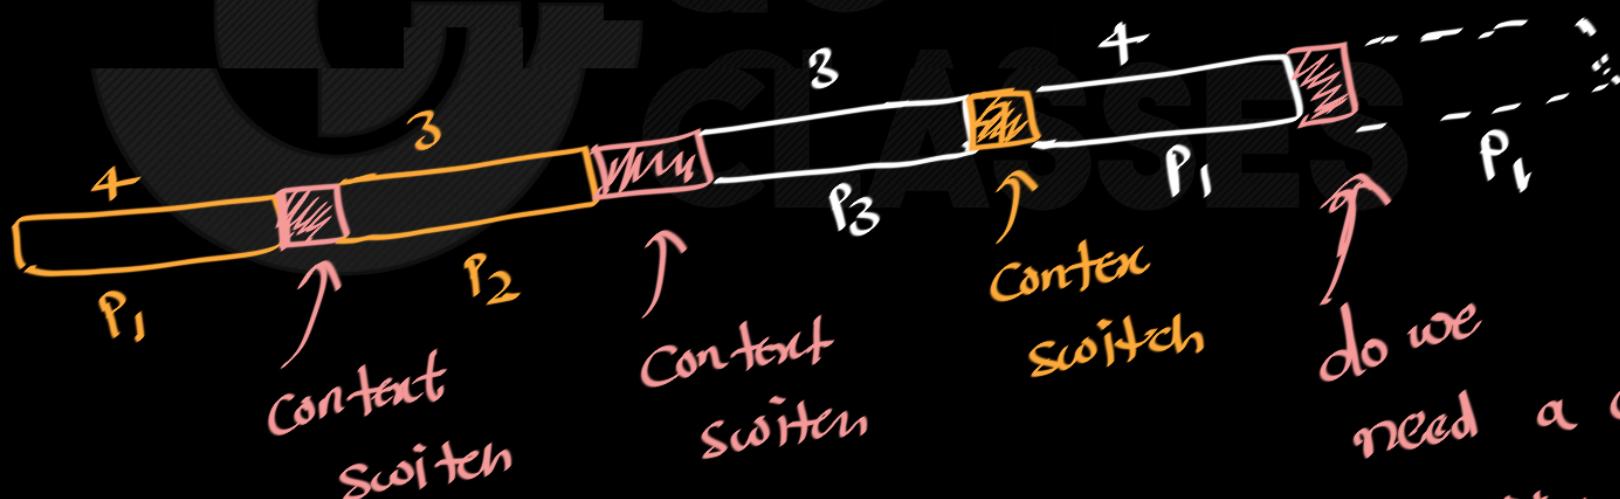
# Operating Systems

Example of RR with Time Quantum = 4



Example of RR with Time Quantum = 4

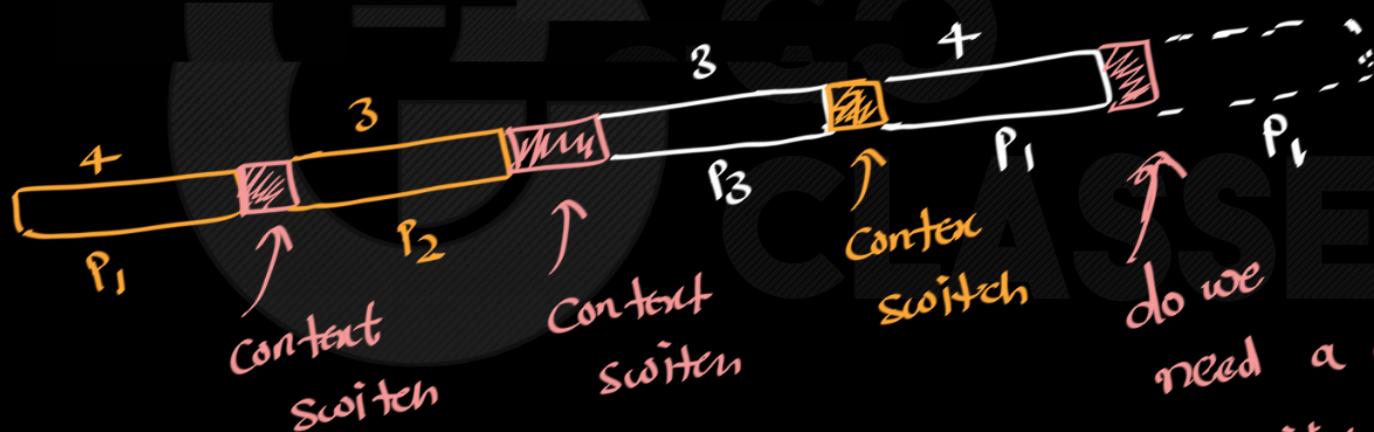
Process	Burst Time	AT	CT	TAT	WT
$P_1$	24	0	30	$30 - 0 = 30$	6
$P_2$	3	0	7	$7 - 0 = 7$	4
$P_3$	3	0	10	$10 - 0 = 10$	7



# Operating Systems

Example of RR with Time Quantum = 4

Process	Burst Time	AT	CT	CT - AT	TAT	WT
$P_1$	24	0	30	$30 - 0 = 30$	30	6
$P_2$	3	0	7	$7 - 0 = 7$	7	4
$P_3$	3	0	10	$10 - 0 = 10$	10	7



Should depend  
on implementation

do we  
need a context  
switch here?

If Ready queue does not have any process then we can bypass (skip) the Context Switch.



# Operating Systems

- In the RR scheduling algorithm, no process is allocated the CPU for more than 1 time quantum in a row (unless it is the only runnable process).



↑  
this won't happen unless  
there is only  $P_1$



## Question

A system uses a Round Robin scheduling algorithm with a time slice of 95 millisecond; the context switch time is 5 milliseconds.

What is the average processor utilization?

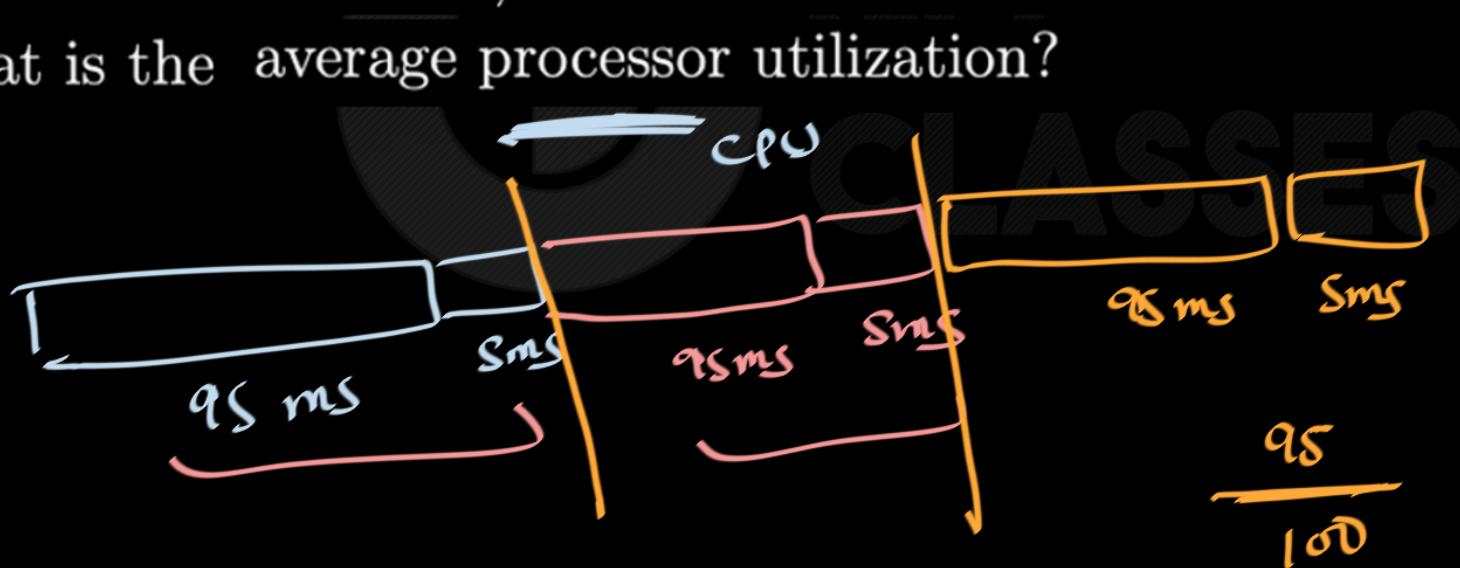


GO CLASSES



## Question

A system uses a Round Robin scheduling algorithm with a time slice of 95 millisecond; the context switch time is 5 milliseconds. What is the average processor utilization?





## Answer

A system uses a Round Robin scheduling algorithm with a time slice of 95 millisecond; the context switch time is 5 milliseconds. What is the average processor utilization?

$$U = \frac{UsefulTime}{TotalTime} = \frac{95}{95+5} = 95\%$$

ASSES  
CPU is working effectively.

$$\text{time quantum} = \delta \text{ ms}$$

$$\text{switch time} = \gamma \text{ ms}$$

$$\text{CPU utilization} = \frac{\text{useful work time}}{\text{total time}} = \frac{\delta}{\delta + \gamma}$$

$$\text{time quantum} = \delta \text{ ms}$$

$$\text{context switch time} = \gamma \text{ ms}$$

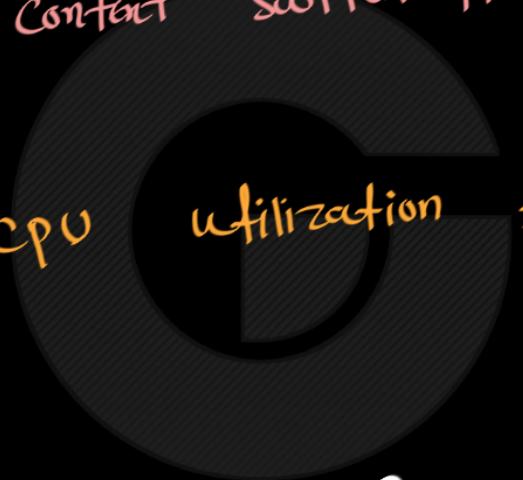
**CPU utilization**

$$\text{CPU utilization} = \frac{\text{useful work time}}{\text{total time}} = \frac{\delta}{\delta + \gamma}$$

$\frac{\delta}{\delta + \gamma} \rightsquigarrow 1$  (will happen when  $\gamma \rightsquigarrow 0$ )

$$\text{time quantum} = \delta \text{ ms}$$

$$\text{context switch time} = \gamma \text{ ms}$$

 CPU utilization =  $\frac{\text{useful work time}}{\text{total time}} = \frac{\delta}{\delta + \gamma}$

Suppose,  $\gamma = \delta$       CPU utilization = 50%.

$$\frac{\delta}{\delta + \gamma} = \frac{\delta}{2\delta} = \frac{1}{2}$$

$$\text{time quantum} = \delta \text{ ms}$$

$$\text{Context switch time} = \gamma \text{ ms}$$

**CPU utilization** =  $\frac{\text{useful work time}}{\text{total time}} = \frac{\delta}{\delta + \gamma}$

$$\frac{\delta}{\delta + \gamma} \rightsquigarrow 1 \quad \left( \begin{array}{l} \text{will happen when} \\ \delta \gg \gamma \end{array} \right)$$

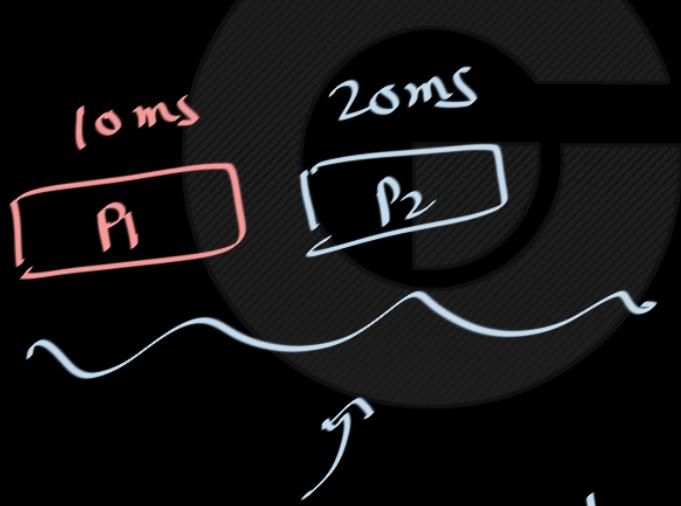
$\gamma$  we can almost ignore  
 $\gamma$  in term  $\delta + \gamma$

$$10000 + 0.001 \approx 10000$$



this is very  
less compare  
to first number (10000)  
then we can ignore this

$$q_f = 10000 \text{ ms}$$



as good as fcfs

for this process

there will be  
no preemption.



- $q$  large  $\Rightarrow$  FIFO
- $q$  small  $\Rightarrow q$  must be large with respect to context switch, otherwise overhead is too high

- Typically, higher average turnaround than SJF, but better **response**
- $q$  should be large compared to context switch time
- $q$  usually 10ms to 100ms, context switch < 10 usec

No mathematical explanation,



# RR advantages and disadvantages

## □ Advantages

- Low response time, good interactivity
- Fair allocation of CPU across processes
- Low average waiting time when job lengths vary widely

## □ Disadvantages

- Poor average waiting time when jobs have similar lengths
  - Average waiting time is even worse than FCFS!
- Performance depends on length of time slice
  - Too high → degenerate to FCFS
  - Too low → too many context switches, costly

SES

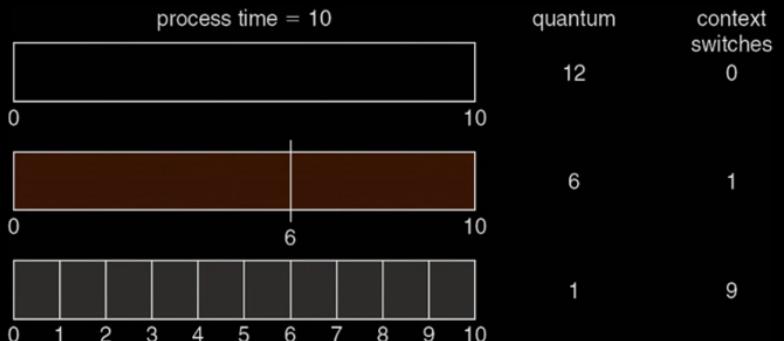


# Round-Robin Scheduling

- RR Scheduling is **preemptive** and designed for time-sharing
- It defines a time quantum
  - A fixed interval of time (10-100ms)
- Unless a process is the only READY process, it never runs for longer than a time quantum before giving control to another ready process
  - It may run for less than the time quantum if its CPU burst is smaller than the time quantum
- Ready Queue is a FIFO
  - Whenever a process changes its state to READY it is placed at the end of the FIFO
- Scheduling:
  - Pick the first process from the ready queue
  - Set a timer to interrupt the process after 1 quantum
  - Dispatch the process

SES

# Picking the Right Quantum



- **Trade-off:**
  - Short quantum: great response/interactivity but high overhead
    - Hopefully not too high if the dispatcher is fast enough
  - Long quantum: poor response/interactivity, but low overhead
    - With a very long time quantum, RR Scheduling becomes FCFS Scheduling
- If context-switching time is 10% of time quantum, then the CPU spends >10% of its time doing context switches
- In practice, %CPU time spent on switching is very low
  - time quantum: 10ms to 100ms
  - context-switching time: 10  $\mu$ s

ES



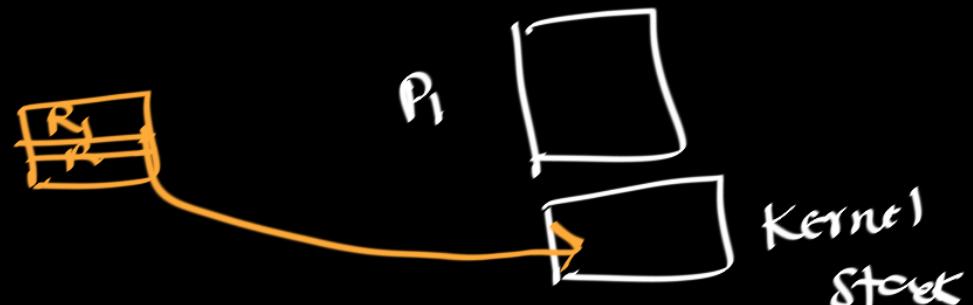
# Operating Systems

## Question

Consider a system with two processes,  $P_1$  and  $P_2$ . Process  $P_1$  is running and process  $P_2$  is ready to run. The operating system uses preemptive round robin scheduling. Consider the situation in which  $P_1$ 's scheduling quantum is about to expire, and  $P_2$  will be selected to run next.

List the sequence of events that will occur in this situation, in the order in which they will occur. Create your list by choosing from the events shown below, using the event numbers to identify events. For example, a valid answer might be: "7,5,4,8,5,2". If an event occurs more than once, you may include it more than once in your list.

1.  $P_1$ 's user-level state is saved into a trap frame
2.  $P_2$ 's user-level state is saved into a trap frame
3. a timer interrupt occurs
4. the operating system's scheduler runs
5. there is a context switch from thread  $T_1$  (from process  $P_1$ ) to thread  $T_2$  (from process  $P_2$ ).
6.  $P_1$ 's user-level state is restored from a trap frame
7.  $P_2$ 's user-level state is restored from a trap frame
8. a "system call" instruction is executed
9. thread  $T_2$  (for process  $P_2$ ) is created
10. thread  $T_1$  (from process  $P_1$ ) is destroyed





# Operating Systems

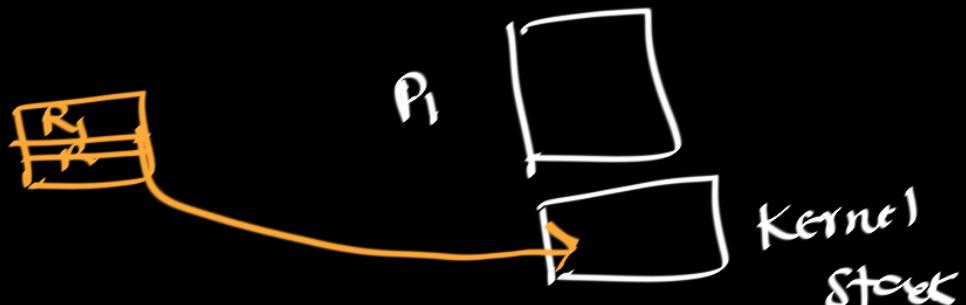
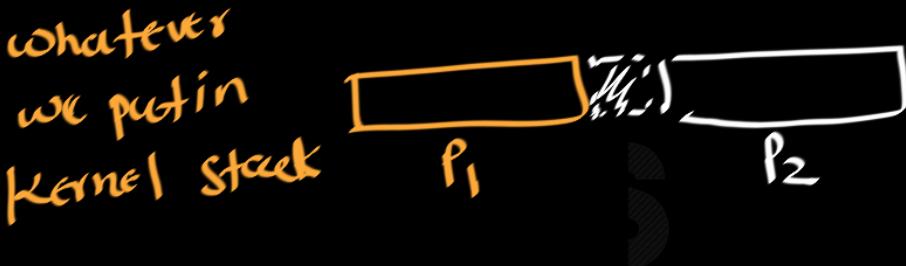
## Question

H.W.

Consider a system with two processes,  $P_1$  and  $P_2$ . Process  $P_1$  is running and process  $P_2$  is ready to run. The operating system uses preemptive round robin scheduling. Consider the situation in which  $P_1$ 's scheduling quantum is about to expire, and  $P_2$  will be selected to run next.

List the sequence of events that will occur in this situation, in the order in which they will occur. Create your list by choosing from the events shown below, using the event numbers to identify events. For example, a valid answer might be: "7,5,4,8,5,2". If an event occurs more than once, you may include it more than once in your list.

1.  $P_1$ 's user-level state is saved into a trap frame
2.  $P_2$ 's user-level state is saved into a trap frame
3. a timer interrupt occurs
4. the operating system's scheduler runs
5. there is a context switch from thread  $T_1$  (from process  $P_1$ ) to thread  $T_2$  (from process  $P_2$ ).
6.  $P_1$ 's user-level state is restored from a trap frame
7.  $P_2$ 's user-level state is restored from a trap frame
8. a "system call" instruction is executed
9. thread  $T_2$  (for process  $P_2$ ) is created
10. thread  $T_1$  (from process  $P_1$  is destroyed





# Operating Systems

Write your answer here:

3,1,4,5,7

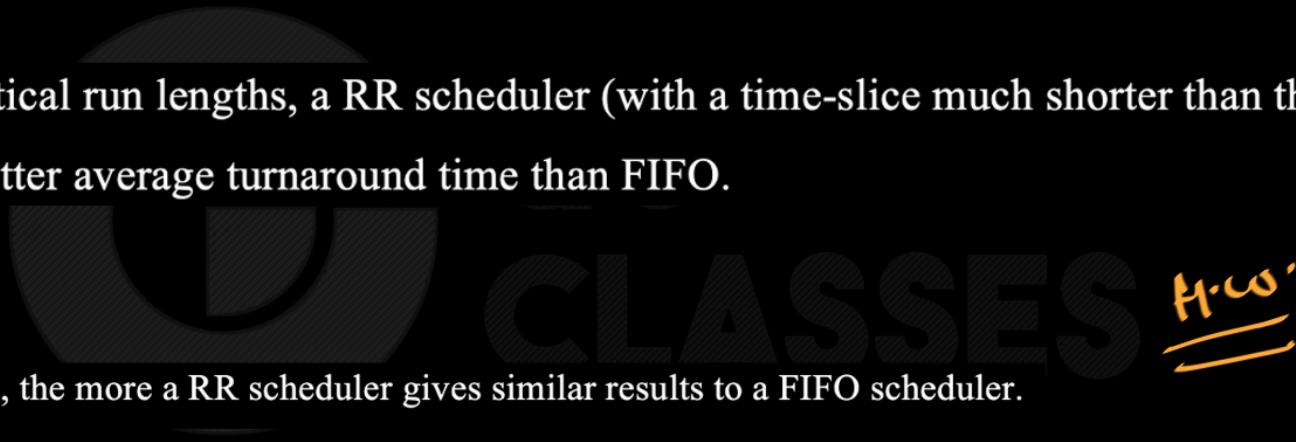




## Question

## True/False

- 1 If all jobs have identical run lengths, a RR scheduler (with a time-slice much shorter than the jobs' run lengths) provides better average turnaround time than FIFO.
  
- 2 The longer the time slice, the more a RR scheduler gives similar results to a FIFO scheduler.



<https://pages.cs.wisc.edu/~dusseau/Classes/CS537/Fall2016/Exams/f16-e2-answers.pdf>



# Operating Systems

- 4) If all jobs have identical run lengths, a RR scheduler (with a time-slice much shorter than the jobs' run lengths) provides better average turnaround time than FIFO.

**False – with RR, identical jobs will all finish at nearly the same time (at the very end of the workload time), which has very poor average turnaround time.**

- 5) The longer the time slice, the more a RR scheduler gives similar results to a FIFO scheduler.

**True – In the extreme, when the time slice is  $\geq$  the length of the job, RR degenerates to FIFO (or FCFS).**

<https://pages.cs.wisc.edu/~dusseau/Classes/CS537/Fall2016/Exams/f16-e2-answers.pdf>



## Question

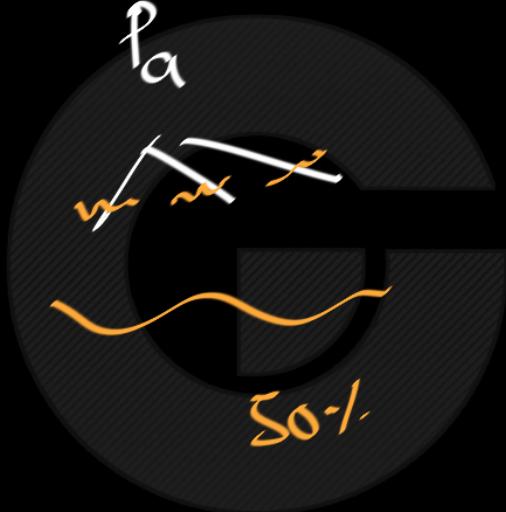


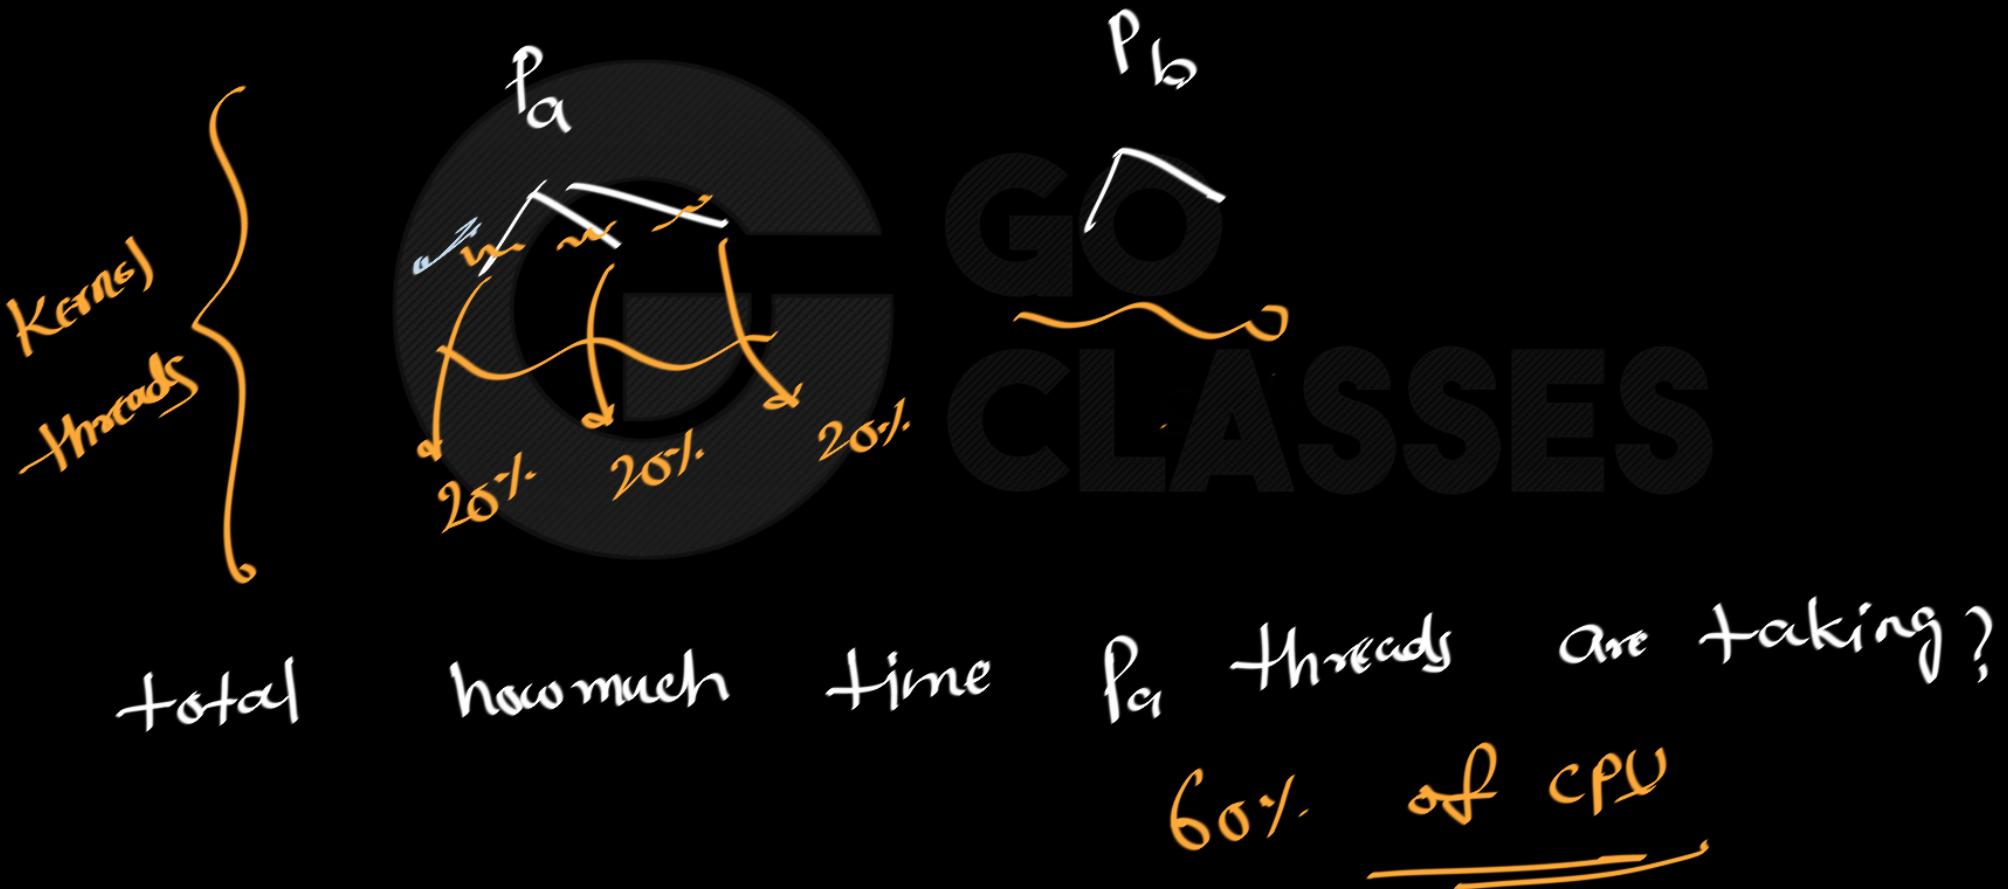
Suppose that two processes,  $P_a$  and  $P_b$ , are running in a uniprocessor system.  $P_a$  has three threads.  $P_b$  has two threads. All threads in both processes are CPU-intensive, i.e., they never block for I/O. The operating system uses simple round-robin scheduling.

- Suppose that all of the threads are user-level threads, and that user-level threads are implemented using a single kernel thread per process. What percentage of the processor's time will be spent running  $P_a$ 's threads? For full credit, briefly justify your answer.
- Suppose instead that all of the threads are kernel threads. What percentage of the processor's time will be spent running  $P_a$ 's threads? For full credit, briefly justify your answer.



user  
threads







# Operating Systems

Suppose that two processes,  $P_a$  and  $P_b$ , are running in a uniprocessor system.  $P_a$  has three threads.  $P_b$  has two threads. All threads in both processes are CPU-intensive, i.e., they never block for I/O. The operating system uses simple round-robin scheduling.

- a. Suppose that all of the threads are user-level threads, and that user-level threads are implemented using a single kernel thread per process. What percentage of the processor's time will be spent running  $P_a$ 's threads? For full credit, briefly justify your answer.

*50% of the processor's time will be spent running  $P_a$ 's threads.*

*Each process has a single kernel thread. Threads never block, so they will always consume their entire scheduling quantum. One a thread has consumed its quantum, it will be preempted in to allow the thread from the other process to run.*

- b. Suppose instead that all of the threads are kernel threads. What percentage of the processor's time will be spent running  $P_a$ 's threads? For full credit, briefly justify your answer.

*60% of the process's time will be spent running  $P_a$ 's threads.*

*In this case there are five kernel threads, each of which will receive 20% of the processor's time. Three of these are associated with process  $P_a$ .*



## ROUND-ROBIN VS. STF/FIFO

- Round-Robin preferable for
  - Interactive applications
  - User needs quick responses from system
- FIFO/STF preferable for Batch applications
  - User submits jobs, goes away, comes back to get result

SES



## Shortest-remaining-time-first (SRTF)

