



Lecture: 18

CLASSES



Operating Systems

```
/* Initialization */  
while (TRUE) {  
    /* entry code */  
    CS() /* critical section */  
    /* exit code */  
    Non_CS()/* non-critical section */  
}
```

No assumption about how often
the critical section is accessed

Wrapper code

ES

Another Try

ME ✓
Progress ✗

int turn = 1 // shared variable

Code for Process 0

while (turn!=0); ← P_0
CS
turn = 1;

initial setup

Code for Process 1

while (turn!=1);
CS
turn = 0;

P_1 P_0 P_1 P_0

BW ✓



Yet Another Try...

1

```
int want[2] = {False, False}; // shared variable
```

Code for Process 0

```
want[0] = True;  
while (want[1]);
```

CS

```
want[0] = False;
```

Code for Process 1

```
want[1] = True;  
while (want[0]);
```

CS

```
want[1] = False;
```

ME ✓
= Progress × (Cause of deadlock)

Program deadlock

921-Omi to Everyone 7:01 PM

9

Sir if shared variable like turn etc
are not initialized so in exam we
have to take all possibilities which
he can take it

↳ yes
if in any of the execution, some condition is not
satisfied then we say overall it is not satisfied.



Peterson's Solution



CLASSES



Peterson's Solution

Code for Process 0

```
want[0] = True;  
favored = 1;  
while (want[1] && favored ==1);
```

CS

```
want[0] = False;
```

```
// shared variables  
int want[2] = {False, False};  
int favored = 0;
```

Code for Process 1

```
want[1] = True;  
favored = 0;  
while (want[0] && favored ==0);
```

CS

```
want[1] = False;
```

{this is having}
deadlock}

Code for Process 0

```
want[0] = True;  
while (want[1] == False);
```

CS

```
want[0] = False;
```

```
// shared variables  
int want[2] = {False, False};  
// Found a bug
```

Code for Process 1

```
want[1] = True;  
while (want[0] == False);
```

CS

```
want[1] = False;
```



Peterson's Solution

it is a polite way
of saying that if a
other process

Code for Process 0

```
want[0] = True;  
favored = 1;  
while (want[1] && favored ==1);
```

CS

```
want[0] = False;
```

i am interested

Code for Process 1

```
want[1] = True;  
favored = 0;  
while (want[0] && favored ==0);
```

CS

```
want[1] = False;
```

```
// shared variables
int want[2] = {False, False};
int favored = 0;
```

Satyam Naik to Everyone 7:16 PM

SN

wow....processes trying to be good
with each other

Rishesh Tiwari to Everyone 7:17 PM

RT

they also have some emotion :)

...

Akash Maji to Everyone 7:17 PM

AM

they have mutual understanding



Purpose of want [] is to tell that, " i am interested"

Purpose of favour is to be polite and say i favour
other process.

"favoured" processes are interested in CS.

will break the tie if both it will break the tie in the by agreeing to the process who executed it last.



Akash Maji to Everyone 7:28 PM

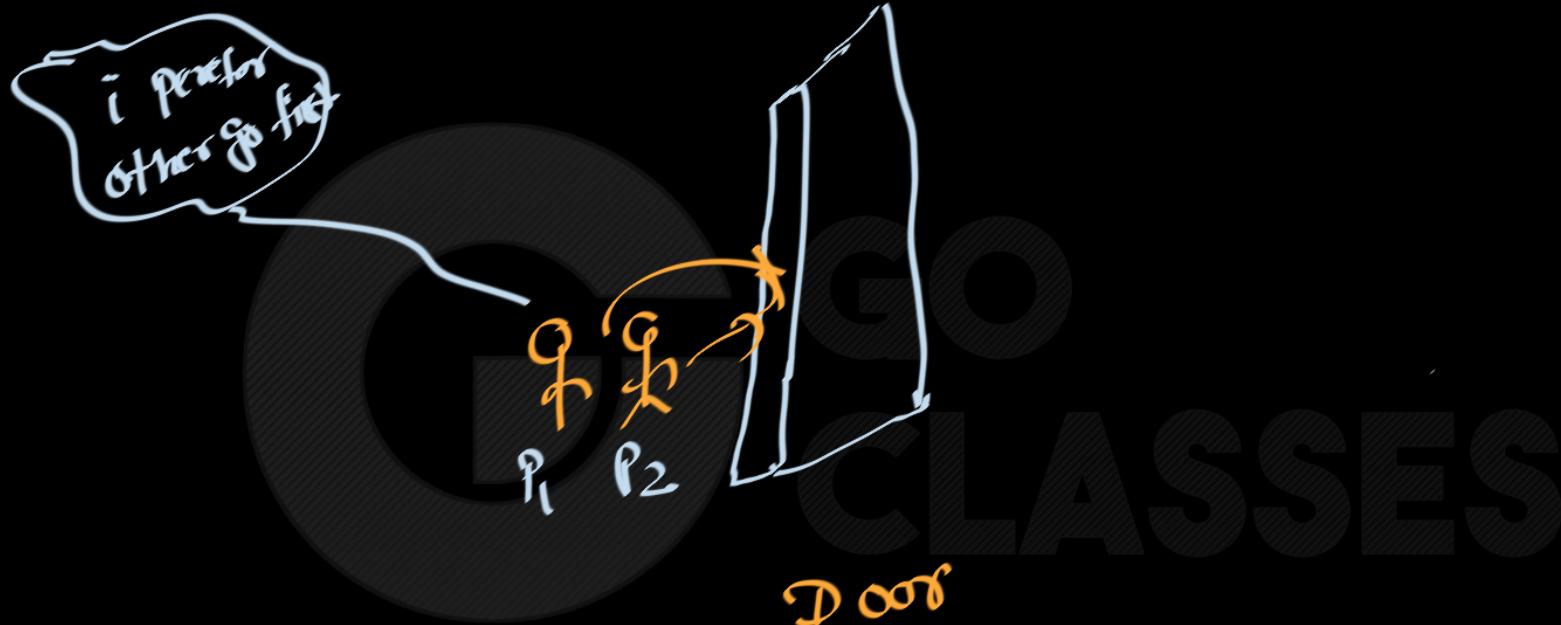
AM

It was very nice way of understanding Peterson solution, having seen similar approaches earlier which did not work and combining them, which is otherwise a bit not so intuitive when directly dealt.

Nishanth D to Everyone 7:30 PM

ND

processes "being polite" helped to understand easily sir



Both the persons interested to go inside.

Me ✓
process ✓

Code for Process 0

```
want[0] = True;  
favored = 1;  
while (want[1] && favored ==1);
```

CS ← P

```
want[0] = False;
```



```
// shared variables  
int want[2] = {False, False};  
int favored = 0;
```

Code for Process 1

```
want[1] = True;  
favored = 0;  
while (want[0] && favored ==0);
```

CS

```
want[1] = False;
```

favored is breaking the tie (in case of deadlock)

Bus ↗

Code for Process 0

```
want[0] = True;
favored = 1;
while (want[1] && favored ==1);
```

CS ← Process 0

```
want[0] = False;
```

Can not
be entered

Code for Process 1

```
want[1] = True;
favored = 0;
while (want[0] && favored ==0);
```

process 1



CS

```
want[1] = False;
```

favored = 0



forward
want

```
int turn;
boolean flag[2];
do {
```

```
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j);
```

critical section

```
    flag[i] = false;
```

remainder section

```
} while (true);
```

ASSES

Source: Galvin

Figure 5.2 The structure of process P_i in Peterson's solution.



```
widget x; //protected variable  
bool she_wants(false);  
bool he_wants(false);  
enum theirs {hers, his} turn;
```

Her Thread

```
she_wants = true;  
turn = his;  
while(he_wants && turn==his);  
frob(x); //critical section  
she_wants = false;
```

His Thread

```
he_wants = true;  
turn = hers;  
while(she_wants && turn==hers);  
borf(x); //critical section  
he_wants = false;
```

*if she wants
and it is her turn
then i will wait*



Operating Systems



Consider the following correct solution to the critical section problem that was discussed in class. The shared variables are the integer turn and the integer array flag[2]. There are two processes, P0 and P1. These processes execute the following code in a loop running on a single CPU with a time sharing operating system.

```
1: P0 code:  
2: while(1) {  
3:   flag[0] = 1;  
4:   turn = 1;  
5:   while((flag[1]==1)&&(turn==1));  
6:     <critical section>  
7:   flag[0] = 0;  
8:     <remainder section>  
9: }
```

```
P1 code:  
while(1) {  
  flag[1] = 1;  
  turn = 0;  
  while((flag[0]==1)&&(turn==0));  
    <critical section>  
  flag[1] = 0;  
    <remainder section>  
}
```

- a. [4pt] What is meant by **bounded waiting**? Does this solution satisfy bounded waiting? If so, what is the bound?
- b. [1pt] What can you say about the values of the shared variables if you know that P1 is in its critical section?
- c. [1pt] What can you say about the values of the shared variables if you know that P1 is in its remainder section?
- d. [2pt] What can you say about the values of the shared variables if you know that P0 is in its critical section and P1 is in its remainder section?
- e. [2pt] What is meant by **busy waiting**? Does this solution use busy waiting?

<http://www.cs.utsa.edu/~korkmaz/teaching/resources-os-ug/tk-samples/solutions/2017c-Fall-cs3733-Final-solution.pdf>

```
1: P0 code:  
2: while(1) {  
3:   flag[0] = 1;  
4:   turn = 1;  
5:   while((flag[1]==1)&&(turn==1));  
6:     <critical section>  
7:   flag[0] = 0;  
8:     <remainder section>  
9: }
```

```
P1 code:  
while(1) {  
  flag[1] = 1;  
  turn = 0;  
  while((flag[0]==1)&&(turn==0));  
    <critical section>  
  flag[1] = 0;  
    <remainder section>  
}
```

turn = , flag []
can be anything

- a. [4pt] What is meant by **bounded waiting**? Does this solution satisfy bounded waiting? If so, what is the bound?
- b. [1pt] What can you say about the values of the shared variables if you know that P1 is in its critical section?
- c. [1pt] What can you say about the values of the shared variables if you know that P1 is in its remainder section?
- d. [2pt] What can you say about the values of the shared variables if you know that P0 is in its critical section and P1 is in its remainder section?
- e. [2pt] What is meant by **busy waiting**? Does this solution use busy waiting?

- c. [1pt] What can you say about the values of the shared variables if you know that P1 is in its remainder section?

```

1: P0 code:
2: while(1) {
3:   flag[0] = 1;
4:   turn = 1;
5:   while((flag[1]==1)&&(turn==1));
6:     <critical section>
7:   flag[0] = 0;
8:     <remainder section>
9: }
```

```

P1 code:
while(1) {
  flag[1] = 1;
  turn = 0;
  while((flag[0]==1)&&(turn==0));
    <critical section>
  flag[1] = 0;
    <remainder section>
}
```



$\nearrow P_1$

turn = could be anything

- d. [2pt] What can you say about the values of the shared variables if you know that P0 is in its critical section and P1 is in its remainder section?

```

1: P0 code:
2: while(1) {
3:   flag[0] = 1;
4:   turn = 1;           ← P0
5:   while((flag[1]==1)&&(turn==1));
6:     <critical section> ← P0
7:   flag[0] = 0;
8:   <remainder section> ← P0
9: }
  
```

```

P1 code:
while(1) {
  flag[1] = 1;
  turn = 0;           ← P1
  while((flag[0]==1)&&(turn==0));
    <critical section> ← P1
  flag[1] = 0;
  <remainder section> ← P1
}
  
```

we want P_1 to go to CS first as per question

$turn = 0$ (P_0 must have executed $\rightarrow P_1$ has execute)

$turn = 1$ (P_1 execute $\rightarrow P_0$ execute)

$turn = 1$



Operating Systems

- a. [4pt] What is meant by **bounded waiting**? Does this solution satisfy bounded waiting? If so, what is the bound?

2pt- Bounded waiting means that once a given process attempts to enter its critical section, there is a bound on the number of times other processes can enter their critical sections before that given process will enter its critical section.

1pt- This solution satisfies bounded waiting

1pt- with a bound of 1.

- b. [1pt] What can you say about the values of the shared variables if you know that P1 is in its critical section?

You only know that **flag[1]** is 1.

Each of the other variables may be either 0 or 1.

- c. [1pt] What can you say about the values of the shared variables if you know that P1 is in its remainder section?

You only know that **flag[1]** is 0.

Each of the other variables may be either 0 or 1.

ASSES



Operating Systems

- d. [2pt] What can you say about the values of the shared variables if you know that P0 is in its critical section and P1 is in its remainder section?

You know that **flag[0]** is 1 and **flag[1]** is 0.

turn must be 1,

otherwise, for **turn** to be 0

P0 must set **turn** to 0 after P0 sets **turn** to 1.

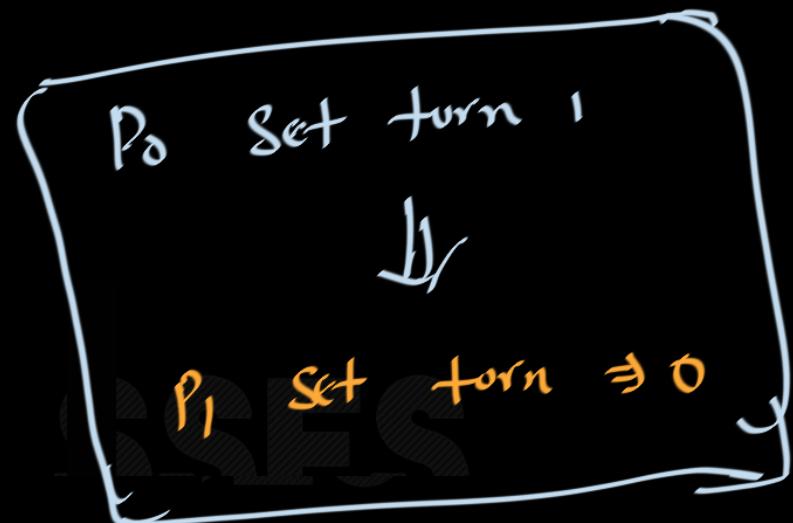
At this time **flag[0]** is 1 and **turn** is 0,

so P1 cannot get past its while loop until P0 finishes its C.S.,
so P1 is not in its remainder section.

- e. [2pt] What is meant by **busy waiting**? Does this solution use busy waiting?

Busy waiting means using the CPU to continually check for an event,

as this solution does in the **while** loop.



P
P1 can not
get in CS first

T_1

while (flag[0]) ; \Rightarrow Spinning

loop till the point

\rightarrow flag[1] = 0 $\left(\begin{array}{l} \text{Busy in waiting} \\ \text{for flag[1] to be} \\ \text{true} \end{array} \right)$

\Rightarrow it is scheduled \Rightarrow taking CPU time and doing nothing



Question :

3. Consider the following solution for the two-process synchronization, where below is the code for P_i (recall that, as usual, P_j is the *other* process where $i = 1 - j$).

```
do {  
    flag[i] = true  
    turn = j  
    while (flag[j] and turn=j);  
  
    CRITICAL SECTION  
  
    flag[i] = false  
  
    REMAINDER SECTION  
  
} while(1)
```

Show that the following three hold:

$j=1$ $j=0$
 $i=0$ $i=1$

- (a) Mutual exclusion is preserved
- (b) Progress requirement is satisfied
- (c) Bounded waiting requirement is met

(Use next blank page if needed.)



Operating Systems

Solution: Mutex: suppose P_0, P_1 are both in the critical section. That would mean that `flag[0]=flag[1]=true` and so they both entered their respective critical sections by breaking out of the while-loop with `turn=0` (for P_0) and `turn=1` (for P_1). But `turn` is a shared variable, and hence this is not possible.

Progress: suppose that P_0 is ready for another round of CS, and that P_1 is in its RS. That means that P_1 just set `flag[1]=false` which breaks P_0 out of the while loop. The dual argument shows the same for P_1 ready for CS and P_0 in its RS.

Bounded waiting: If P_0 is ready to enter its CS, then P_1 will be able to enter its CS at most once; as soon as P_1 exits its CS, it turns `flag[1]=false` allowing P_0 entry. Again, the dual argument shows the same for P_1 and P_0 .



We have seen software based solutions, which means all the sol'n's were just piece of code.

Do you think any disadvantage of Peterson's sol'n.

Disadvantage of Peterson's solⁿ (or any other working solⁿ)

- Code are slow (compared to hardware)
- Proving the correctness is also an effort
- Peterson's solⁿ is only for 2 processes
 - we can generalise to n process but n must be known in prior. (out of syllabus)
- Busy waiting (wasting CPU cycles)



Hardware Supported Solutions

GO
CLASSES



Hardware Supported Solutions

- Disabling Interrupt
- Special Machine Instructions
 - test-and-set, compare-and-swap, ...





1. Hardware Solutions: Disabling Interrupt

- The critical-section problem could be solved simply in a single-processor environment if we could prevent interrupts from occurring while a shared variable was being modified.

```
while (true) {  
    /* disable interrupts */;  
    /* critical section */;  
    /* enable interrupts */;  
    /* remainder */;  
}
```

if you disable the
interrupts before entering
into CS, it will disable
any kind of interrupt
including context switch

in CS, CPU will just be executing
CS without preemption



Working of Disabling Interrupt

The critical-section problem could be solved simply in a single-processor environment if we could prevent interrupts from occurring while a shared variable was being modified. In this way, we could be sure that the current sequence of instructions would be allowed to execute in order without pre-emption. No other instructions would be run, so no unexpected modifications could be made to the shared variable.

Source: Galvin





Problems with Disabling Interrupt

Unfortunately, this solution is not as feasible in a multiprocessor environment.

→ for every process, we have dedicated "interrupt lines"

Source: Galvin



Problems with Disabling Interrupt

Unfortunately, this solution is not as feasible in a multiprocessor environment.

- for every process, we have dedicated “interrupt lines”
- 
- P₂ (P₂ is dealing with same CS as P₁ is dealing)

Source: Galvin



Problems with Disabling Interrupt

Unfortunately, this solution is not as feasible in a multiprocessor environment. Disabling interrupts on a multiprocessor can be time consuming, since the message is passed to all the processors. This message passing delays entry into each critical section, and system efficiency decreases. Also consider the effect on a system's clock if the clock is kept updated by interrupts.

Source: Galvin



Problems with Disabling Interrupt (cont..)

This approach is generally unattractive because it is unwise to give user processes the power to turn off interrupts. What if one of them did it, and never turned them on again? That could be the end of the system. Furthermore, if the system is a multiprocessor (with two or more CPUs) disabling interrupts affects only the CPU that executed the **disable** instruction. The other ones will continue running and can access the shared memory.

Source: Tanenbaum

Interrupt Disabling

In a uniprocessor system, concurrent processes cannot have overlapped execution; they can only be interleaved. Furthermore, a process will continue to run until it invokes an OS service or until it is interrupted. Therefore, to guarantee mutual exclusion, it is sufficient to prevent a process from being interrupted. This capability can be provided in the form of primitives defined by the OS kernel for disabling and enabling interrupts. A process can then enforce mutual exclusion in the following way (compare to Figure 5.4):

```
while (true) {  
    /* disable interrupts */;  
    /* critical section */;  
    /* enable interrupts */;  
    /* remainder */;  
}
```

Because the critical section cannot be interrupted, mutual exclusion is guaranteed. The price of this approach, however, is high. The efficiency of execution could be noticeably degraded because the processor is limited in its ability to interleave processes. Another problem is that this approach will not work in a multiprocessor architecture. When the computer includes more than one processor, it is possible (and typical) for more than one process to be executing at a time. In this case, disabled interrupts do not guarantee mutual exclusion.

Optional read from
William Stalling

Other solⁿ that we seek from

hardware : H/w itself provide

atomic "Read - Modify - Write" instructions

(RMW)





Rishesh Tiwari to Everyone 8:27 PM

RT

You done post - mortam of
everything of Peterson solution
today :)



1

GO
CLASSES



2. Hardware Solutions: Atomic operations

CLASSES



Atomic operations

Processors provide means to execute **read-modify-write** operations atomically on a memory location

- Typically applies to at most 8-bytes-long variables

Common atomic operations

- `test_and_set(type *ptr)`: sets `*ptr` to 1 and returns its previous value
- `fetch_and_add(type *ptr, type val)`: adds `val` to `*ptr` and returns its previous value
- `compare_and_swap(type *ptr, type oldval, type newval)`: if `*ptr == oldval`, set `*ptr` to `newval` and returns true; returns false otherwise

SES



- **Examples:**

- **Test&Set:** (most architectures) read a value, write ‘1’ back to memory.
- **Exchange:** (x86) swaps value between register and memory.
- **Compare&Swap:** (68000) read value, if value matches register value r1, exchange register r2 and value.



<https://lass.cs.umass.edu/~shenoy/courses/spring10/lectures/Lec08.pdf>